

Waterfall IR Design

Research notes + IR reference

BMA Standard Formulas

May 02, 2026

Waterfall IR Design — research notes + IR reference

Methodology

RMBS — Agency MBS REMIC

FNR 2006-018 (existing fixture, anchor)

FNR 2016-104 (Dec 2016)

FNR 2019-17 (Mar 2019)

Common patterns across all 3 agency REMIC deals

Agency Synthetic CRT (CAS 2024-R05, CAS 2024-R06)

Structure

Cashflow mechanics

What CRT needs from the IR

Match with current IR

RMBS — Non-Agency (Subprime, JPMMT 2006)

Sample read: JPMorgan Mortgage Acceptance 2006 (CWHEQ-style)

Auto ABS — Prime (Ford Credit Auto Owner Trust 2024-C)

Pre-acceleration priority of payments (single 11-step cascade)

Key structural features

What's notable vs RMBS

Auto ABS — Subprime (Santander Drive Auto Receivables Trust 2024-2)

Pre-acceleration priority of payments (single 14-step cascade)

Same shape as prime auto with parametric differences only

Post-acceleration priority of payments

Cross-asset-class observations

What ALL deals have in common

What VARIES across asset classes

Cross-cutting needs

Gaps in the current IR

Proposed IR additions (DRAFT — not approved)

A. Multi-target rule consolidation (no schema change, fixture rewrite + synth/irGen change)

B. ComputedAmount node (small schema add)

C. WaterfallBranch — `if / elif / else` over rule blocks

D. AggregateGroup bond bundle

E. Bond-level loss treatment (the real loss-allocation question)

F. Recursive SPLIT_CASH (no schema change, runtime change)

Round 3 schema review — separating bond identity from rule behavior

G. Collapse TrancheType + TrancheBehavior into a single TrancheKind

H. Unify tranche relationships into one structured list

I. Coupon as a value or a schedule (not just a scalar)

J. Notional, not "size"

K. Two-phase schedule derivation — speeds are design-time, `schedule_contract` is runtime

L. Drop CONCURRENT from PaymentStyle

M. Reserve rules are not a separate rule type; they're rules sourced from / targeted to accounts

N. Drop alias tokens; rename INT_CASH / PRIN_CASH

O. AccountType is a label, not runtime semantics

Q. Implement `minimum_basis` and `starting_basis` per-period semantics — FIXED

P. Schema cleanup migration table

Open questions for user before any code change

Additional deals (round 2 research)

Ginnie Mae REMIC 2025-203 (single-family, multi-group)

Ginnie Mae REMIC 2025-009 (HECM-backed reverse mortgage)

Ginnie Mae REMIC 2024-115 (Multifamily)

Freddie Mac REMIC general structure (offering circular)

Verus Securitization Trust 2024-9 (Modern Non-QM RMBS)

Toyota Auto Receivables 2024-A (Prime Auto, Toyota)

Toyota Lexus Owner Trust 2024-A (Auto LEASE — different shape)

Westlake Automobile Receivables Trust 2024-1 (Subprime Auto)

Updated cross-asset matrix

Updated IR gap list

Part 2 — IR reference: the schema and how to write a deal

IR overview

Top-level schema: `DealDefinition`

`deal_knobs` — what's actually in there

Migration intent for `deal_knobs`

`BondDef` — every tranche the deal pays

`RuleNode` — one waterfall step

`RuleType` enum — what cash this rule moves

`PaymentStyle` — order semantics within a multi-target rule

`CapMode` — schedule cap interpretation for PAY_PRINCIPAL

Built-in source/target tokens (reusable across asset classes)

Other elements

`AccountDef` — named cash buckets (reserves, prefunding, etc.)

`FeeDef` — periodic fee paid to a payee

`TriggerNode` — conditional gate

`CalculationNode` — named expressions

`CollateralGroupDef` — multi-pool deals

Worked example: FNR 2006-018 Group 1, written in IR

Reusability principles

Human-readability principles

Anti-patterns to avoid

Waterfall IR Design — research notes + IR reference

Status: Initial design synthesis from 13 prospectus deep reads across RMBS and auto ABS, plus full IR element reference and a worked example. Not a final spec for IR change yet. Part 1 documents what real prospectus language looks like; Part 2 documents the existing IR schema and how to write a deal in it both reusably and readably.

Sample size (13 deals + FNR 2006-018 fixture):

Asset class	Deal	Key features
Agency MBS REMIC	FNR 2006-018 (fixture)	2 collateral groups, PAC + Z + Support, face-weighted support split
Agency MBS REMIC	FNR 2016-104	9 collateral groups, mix of pass-through, sequential, accretion-directed, PAC, face-weighted splits
Agency MBS REMIC	FNR 2019-17	7 collateral groups, nested face-weighted splits , named Aggregate Group abstraction
Agency Multifamily REMIC	FNMA 2024-M2	Multifamily; structurally similar to single-family REMICs
Agency Synthetic CRT	CAS 2024-R05, CAS 2024-R06	Connecticut Avenue Securities — synthetic credit risk transfer; reference pool of FNMA-acquired loans, M-1 / M-2 / B-1 / B-2 notes, reverse-seniority bond writedowns on reference-pool losses (sometimes with later writeup), pro-rata principal pre-stepdown / sequential post-stepdown
Agency MBS REMIC	Ginnie Mae 2025-203	Confirms the FNR PAC + Z + Support pattern is industry-standard ; “Aggregate Scheduled Principal Balance” same abstraction as Fannie’s “Aggregate Group Planned Balance”
Agency MBS REMIC	Ginnie Mae 2025-009 (HECM)	Reverse-mortgage REMIC; Deferred Interest Amount (catch-up rule type not in current IR)
Agency MBS REMIC	Ginnie Mae 2024-115 (Multifamily)	Multifamily-specific: trustee fee % of Principal Distribution Amount before cascade
Freddie Mac REMIC general OC	(offering circular)	Single-Tier vs Double-Tier Series : REMIC-inside-REMIC trust structure (mostly transparent for cashflow IR)
Non-Agency RMBS (subprime)	JPMMT 2006	Single pool, interest waterfall + principal waterfall sub-streams, stepdown date , trigger event override , OC + excess interest, M-1..M-10 mezz with reverse-seniority loss allocation
Non-Agency RMBS (Non-QM)	Verus 2024-9 / 2026-4	Step-up coupon at year 5 (time-conditional rate), LCF (Last Cashflow) class , mixed pay (senior pro rata, mezz/sub sequential), Class XS for excess spread
Prime Auto ABS	Ford Credit Auto Owner Trust 2024-C	Single pool, interleaved I/P , named priority principal amounts, target OC build, reserve replenishment, capped trustee fee with overflow
Prime Auto ABS	Toyota Auto Receivables 2024-A	Same shape as Ford Credit. Yield Supplement Overcollateralization Amount for sub-WAC loan adjustment. Confirms prime auto = same grammar across sponsors.
Auto Lease ABS	Toyota Lexus Owner Trust 2024-A (TLOT)	Lease-specific simpler waterfall (8 steps): pro-rata interest across all classes, single combined priority principal amount, “Securitization Value” valuation. Otherwise same grammar.
Subprime Auto ABS	Santander Drive 2024-2 (SDART)	Same shape as Ford / Toyota prime auto . Parametric differences (4 mezz classes vs 2-3, 5 named allocations vs 3, \$300K trustee cap vs \$375K).
Subprime Auto ABS	Westlake 2024-1 (WLAKE)	8-class structure (A-1, A-2, A-3 + B, C, D, E). Same waterfall shape as SDART. Confirms subprime auto pattern is consistent across sponsors.

Key cross-asset finding: every asset class reduces to a small set of structural primitives. Prime auto + subprime auto + auto lease all share one grammar. Agency MBS REMICs across Fannie / Freddie / Ginnie share the same PAC + Z + Support pattern. The existing IR captures most of these patterns; the major real gap is **named computed distribution amounts** (heavy in non-agency RMBS and auto), and the **authoring practice** of fragmenting waterfall steps into one rule per bond (a fixture / generator issue, not an IR limitation).

Asset classes still to cover (not yet in sample): credit card master trusts, CLOs (managed), marketplace consumer (SoFi / LendingClub / Affirm), equipment ABS, aircraft ABS, solar ABS.

Methodology

For each deal: locate the “Distributions” / “Priority of Payments” section verbatim, categorize each numbered step in the prospectus, and note what abstractions the prospectus implicitly assumes (named amounts, computed amounts, named groups).

RMBS — Agency MBS REMIC

FNR 2006-018 (existing fixture, anchor)

Two collateral groups, Group 1 has PAC + Z + Support classes with a 95.65 / 4.35 face-weighted support split. Group 2 is a sequential cascade with notional IO.

FNR 2016-104 (Dec 2016)

9 separate collateral groups in one trust. Each group has its own mini-waterfall. Sample group structures:

- **Group 1 (Pass-Through):** “Group 1 Principal Distribution Amount to BA until retired.”
 - Single rule, single target.
- **Group 4 (Z + Sequential):** Two rules:
 1. “The Z Accrual Amount to A and B, **in that order**, until retired, and **thereafter** to Z.”
 2. “The Group 4 Cash Flow Distribution Amount to A, B and Z, **in that order**, until retired.”
 - Each rule has *multiple* targets in a sequential list.
- **Group 5 (PAC + Support):** Three numbered phases:
 1. To Aggregate Group **to its Planned Balance**.
 2. To LZ until retired.
 3. To Aggregate Group **to zero**.
 - Phase 1 = PAC schedule cap (`cap_mode=PLANNED`).
 - Phase 2 = pure sequential, no cap.
 - Phase 3 = cleanup (`cap_mode=NONE`).
 - **The whole “PAC + Support” pattern is three numbered rules, not three rules per bond.**
- **Group 6 (Face-Weighted Split):**
 - “The Group 6 Cash Flow Distribution Amount as follows:
 - 67.8420172533% to QA, QV and QZ, in that order, until retired
 - 32.1579827467% to QB, QW and QY, in that order, until retired”
 - Single SPLIT_CASH at the named percentage with two sub-cascades.

FNR 2019-17 (Mar 2019)

7 collateral groups. Adds two patterns we did not see in 2006-018 or 2016-104:

- **Pro-rata pay** as a top-level option: “Group 1 Principal Distribution Amount to A and FA, **pro rata**, until retired.” — `payment_style=PRO_RATA` on a multi-target rule.
- **Nested face-weighted splits** in Group 7:
 - “Group 7 ... Aggregate Group III as follows:
 - 16.6666666667% to MA until retired, **and**
 - 83.3333333333% as follows:
 - first, 44.8738331794% to ME until retired, **and** 55.1261668206% to MG and MB, in that order, until retired,
 - second, to MC until retired.”
 - The 83% branch is itself a SPLIT_CASH with two sub-streams, AND that split has two sequential phases (“first” / “second”).
 - **Recursive structure:** rules can contain sub-rules.

- **Named “Aggregate Group” abstraction:** “Aggregate Group III consists of the MA, ME, MC, MG and MB Classes ... has a principal balance equal to the aggregate principal balance of the Classes included in Aggregate Group III.”
 - A bond *bundle* with its own planned balance and internal allocation rules. **Not a separate REMIC class** — a virtual grouping for schedule cap and downstream rules.

Common patterns across all 3 agency REMIC deals

- One waterfall per collateral group.
- Each waterfall has 1-3 logical steps.
- **Each step is a multi-target rule**, not one rule per bond.
- Order language is verbatim from the prospectus: “in that order, until retired”, “concurrently / pro rata”, “to its Planned Balance”, “to zero”.
- Z-bond accrual is its own pre-step: “Accrual Amount to X until retired, then to Z” (a sequential cascade itself).
- Face-weighted splits are first-class and **can nest**.
- “Aggregate Group” is a named bond bundle with its own schedule.

Agency Synthetic CRT (CAS 2024-R05, CAS 2024-R06)

Connecticut Avenue Securities (CAS) — Fannie Mae’s flagship credit risk transfer program. Structurally distinct from cash REMICs but **fully in scope** for the IR.

Structure

- **Reference pool**, not a real cash pool. The trust holds no mortgage collateral — it holds Fannie Mae’s payment obligations derived from the *performance* of a reference pool of FNMA-acquired loans. The reference pool’s amortization, prepayments, and credit events drive the bonds.
- **Note classes** (typical CAS structure): M-1, M-2, B-1, B-2 plus unrated B-3H. M-1 is most senior of the issued notes; B-2 is most junior. There is also a hypothetical reference tranche stack (A-H, M-1H, M-2H, B-1H, B-2H, B-3H) used to compute payments but with no actual notes.
- **No reserves, no excess spread, no OC.** Credit enhancement is pure subordination — junior notes absorb losses before senior.

Cashflow mechanics

Two simple things happen each period:

1. **Interest** — each note accrues coupon (typically SOFR + spread) on its outstanding balance and Fannie Mae pays that coupon monthly. Like a normal floater.
2. **Principal & writedowns** — the reference pool’s scheduled amortization + prepayments are allocated to bonds (pro-rata pre-stepdown, sequential post-stepdown), and any reference-pool credit losses are *written down* against bond balances in reverse seniority (B-2 first, then B-1, M-2, M-1).

What CRT needs from the IR

CRT is the **simplest waterfall structurally** but the **hardest without `BondDef.loss_treatment`**. Because there is no cash collateral, the bonds’ economic existence IS their balance — and losses are the primary state change. Specifically:

- **Bond writedown semantics** — when LOSS is allocated to a bond, its balance must decrease and future coupon must accrue on the reduced balance. This is `BondDef.loss_treatment = WRITEDOWN` (proposed addition E in this document).
- **Bond writeup semantics** (rare but real) — if Fannie Mae later determines a credit event was overstated, the writedown is reversed. The bond’s balance increases and the bondholder receives a “writeup payment” representing missed coupon. This is `BondDef.writeup_enabled = true`.
- **Stepdown date + performance triggers** — same idea as non-agency RMBS: pre-stepdown is pro-rata, post-stepdown is sequential, but a delinquency trigger reverts to sequential early if collateral underperforms. Expressible with the proposed `WaterfallBranch` (proposed addition C).
- **Reference-pool collateral input** — the IR’s existing `Loan` / `DealRunInput` types accept the reference pool’s cashflows directly (treated as if they were real); this part works today.

Match with current IR

CRT feature	Current IR	Status
Sequential / pro-rata principal cascade	<code>PAY_PRINCIPAL</code> with <code>payment_style</code>	✓
Reverse-seniority loss cascade	<code>PAY_WRITEDOWN</code> with reverse <code>to_targets</code>	✓
Bond balance writedown on loss	none	✗ (need <code>loss_treatment</code>)
Bond writeup on loss reversal	none	✗ (need <code>writeup_enabled</code> + <code>PAY_WRITEUP</code>)

CRT feature	Current IR	Status
Stepdown × trigger conditional waterfall	per-rule <code>condition_trigger</code>	⚠️ works but verbose; <code>WaterfallBranch</code> is cleaner
Floater coupon	<code>BondDef.coupon_type=FLOATING</code> + index	✅

The summary: **CRT is structurally the simplest deal in this research corpus, but the existing IR cannot model it correctly because of the missing bond-level loss treatment.** Fixing that one gap unlocks the entire CRT product family (CAS, STACR, CIRT, ACIS).

RMBS — Non-Agency (Subprime, JPMMT 2006)

Sample read: JPMorgan Mortgage Acceptance 2006 (CWHEQ-style)

This deal is structurally **very different** from agency:

- Two collateral groups (Group 1 / Group 2 mortgage loans).
- **Two parallel waterfalls per period** — Interest Remittance Amount cascade and Principal Remittance Amount cascade — that intersect via OC mechanics.
- 12-step **interest waterfall** plus 5-condition **principal waterfall** (varies by stepdown date and trigger event).
- M-1 through M-10 subordinate stack.
- Excess interest cascade (“Net Monthly Excess Cashflow”) with its own multi-step waterfall.
- Net WAC reserve fund with its own sub-waterfall.

Interest waterfall (12 numbered steps)

1. Senior interest (concurrent / pro-rata across A-1A, A-1B, A-2..A-5)
2. Senior unpaid interest shortfalls (concurrent) 3-12. M-1 through M-10 interest, **one step per class**, in strict seniority order.

The prospectus phrases steps 3-12 as one step per mezz class — **NOT** as one rule with [M-1..M-10]. The reason: each step is its own waterfall priority where remaining funds at that step depend on what was paid above. (For interest where shortfalls are tracked per-class, step-by-step is more readable.) **Either rendering (per-class steps OR a single multi-target sequential rule) produces identical math** — but the prospectus authors prefer per-class for clarity in the mezz stack.

Principal waterfall — gated by 2 conditions, 6 conditional blocks

Each block applies under a different **(stepdown date, trigger event)** combination:

- **A:** prior to stepdown OR Trigger Event in effect → Group 1 Principal → Group 1 Certs first, then Group 2 Certs.
- **B:** prior to stepdown OR Trigger Event in effect → Group 2 Principal → Group 2 Certs first, then Group 1 Certs.
- **C:** prior to stepdown OR Trigger Event in effect → remaining combined principal → M-1, M-2, ..., M-10 sequentially.
- **D:** post-stepdown AND no Trigger Event → Group 1 Senior Principal Distribution Amount (a different formula) → Group 1 Senior + cross-allocation to Group 2.
- **E:** post-stepdown AND no Trigger Event → Group 2 Senior Principal Distribution Amount → Group 2 Senior + cross-allocation.
- **F:** post-stepdown AND no Trigger Event → remaining → mezz with **Combined Class M-1, M-2 and M-3 Principal Distribution Amount** (a SHARED computed amount that pays all three sequentially).

Computed distribution amounts

Heavy use of NAMED FORMULAS:

- “Group 1 Principal Distribution Amount” = pool collections + advances
- “Group 1 Senior Principal Distribution Amount” = capped formula
- “Combined Class M-1, M-2 and M-3 Principal Distribution Amount” = formula

These are not rules — they are **named scalar calculations** computed each period and referenced by rules. The current IR has `CalculationNode` for trigger metrics but doesn’t use it for distribution amounts.

Sequential Trigger Event vs Trigger Event

There are **multiple distinct triggers** controlling different behaviors:

- “Trigger Event” — gates the stepdown.
- “Sequential Trigger Event” — gates Class A-1A / A-1B concurrent-vs-sequential.

Both are computed each period from cumulative loss + delinquency ratios with **time-dependent thresholds** (a different threshold table per Distribution Date range).

Loss allocation — separate cascade

Loss allocation is a SEPARATE waterfall: realized losses go first to M-10 → M-9 → ... → M-1 → seniors. **Reverse seniority order**, not the cash distribution order. Currently the IR has no first-class loss-allocation primitive; it's bolted onto bond writedown logic.

Excess interest cascade

"Net Monthly Excess Cashflow" — what's left after paying interest + fees + scheduled principal — has its OWN multi-step waterfall:

1. To OC build until target reached
2. To unpaid interest shortfalls
3. To realized losses (write-up M-10..M-1)
4. To net WAC carryover
5. ... and several more

This is a sub-waterfall on a derived stream. The IR's **SPLIT_CASH** primitive can route to a "Net Monthly Excess Cashflow" stream, and then a sequence of rules consumes it. Workable, but requires expressing the cascade explicitly.

Auto ABS — Prime (Ford Credit Auto Owner Trust 2024-C)

Pre-acceleration priority of payments (single 11-step cascade)

1. Transaction Fees and Expenses (capped \$375K/yr) — to indenture trustee, owner trustee, ARR
2. Servicing Fee — to servicer
3. Class A Note Interest — pro rata across Class A notes
4. First Priority Principal Payment — to Class A, sequentially by class, amount = MAX(0, Class A principal - adjusted pool balance)
5. Class B Note Interest
6. Second Priority Principal Payment — to Class A+B, sequentially by class, amount = MAX(0, (A+B principal) - adjusted pool balance) - step 4 amount
7. Class C Note Interest
8. Reserve Account Replenishment — to top up to original balance
9. Regular Principal Payment — sequentially by class, amount = GREATER OF (a) Class A-1 principal, (b) notes principal - (pool balance - target OC) reduced by first + second priority amounts
10. Additional Fees and Expenses (uncapped overflow from step 1)
11. Residual Interest — to depositor

Key structural features

1. **Single Available Funds source** — no group split. All collections pool together; the waterfall consumes them in 11 steps.
2. **Interest and principal interleaved** by class seniority, not as separate sub-waterfalls. Step 3 = A interest. Step 5 = B interest. Step 7 = C interest. The 4 / 6 / 9 between are principal acceleration / catch-up steps.
3. **Computed principal amounts:**
 - "First Priority Principal Payment" — explicit formula based on overcollateralization deficit; only nonzero if Class A principal exceeds adjusted pool balance.
 - "Second Priority Principal Payment" — same shape, A+B view.
 - "Regular Principal Payment" — explicit MAX-of-two-formulas with reductions for upstream payments.
4. **Reserve account replenishment as a numbered step** (step 8). The current IR has **PAY_TO_RESERVE** rule type, so this is covered, but the name is hidden from the user in the FNR fixture (no reserve in agency).
5. **Sequential by class for principal**, but **pro rata within class** for interest. Class A has sub-classes A-1, A-2a, A-2b, A-3, A-4; principal pays A-1 first to retired, then A-2a, etc.; interest pays all 5 pro rata each period.
6. **Capped fees with overflow:** Step 1 caps at \$375K/year; any excess fees move to step 10 (after principal). The IR's **max_amount_fixed** covers the cap; the carry-to-later-step needs thought.
7. **Post-acceleration priority of payments:** a SEPARATE waterfall for after Event of Default (not detailed in our extract). Adds an ALTERNATE TOP-LEVEL waterfall gated by deal state.

What's notable vs RMBS

- **No PAC / TAC / Z.** Auto deals are pure sequential.
- **Single pool, no groups.**
- **No stepdown gate** the way RMBS has it; instead the auto structure relies on the priority principal mechanism (steps 4 + 6) to maintain seniority backstop, and the "Regular Principal" formula (step 9) to build target OC over time.
- **Reserve account is a first-class participant** in the waterfall, not just a credit enhancement footnote.

- **Asset-rep fees** (a relatively recent addition tied to ABS Rule 15Ga-1) appear as a discrete payee.

Auto ABS — Subprime (Santander Drive Auto Receivables Trust 2024-2)

Pre-acceleration priority of payments (single 14-step cascade)

1. first — trustee + Delaware trustee + owner trustee + ARR fees (\$300K/yr cap)
2. second — servicing fee
3. third — Class A interest, pro rata
4. fourth — First Allocation of Principal
5. fifth — Class B interest
6. sixth — Second Allocation of Principal
7. seventh — Class C interest
8. eighth — Third Allocation of Principal
9. ninth — Class D interest
10. tenth — Fourth Allocation of Principal
11. eleventh — reserve replenishment (to Specified Reserve Balance)
12. twelfth — Regular Allocation of Principal
13. thirteenth — trustee fee overflow (per-annum cap exceeded)
14. fourteenth — residual to certificateholders, pro rata

Same shape as prime auto with parametric differences only

The waterfall is structurally **identical** to Ford Credit prime auto. The differences are entirely parametric:

Feature	Ford Credit (prime)	SDART (subprime)
Step count	11	14
Mezz classes	A, B, C	A, B, C, D
Named principal allocations	First, Second, Regular	First, Second, Third, Fourth, Regular
Trustee fee per-annum cap	\$375K	\$300K
Residual recipient	“Holder of residual interest”	“Certificateholders, pro rata”

Implication for the IR: prime + subprime auto share **one waterfall grammar**. The IR does not need a “subprime auto” rule type or a “prime auto” rule type — same primitives, more parameters.

Post-acceleration priority of payments

After Event of Default + acceleration, a SEPARATE waterfall applies (typically: trustee fees → all classes interest pro rata → all classes principal sequentially with no priority principal acceleration). Triggered by indenture conditions, not period-by-period.

This is structurally the same as the JPMMT “Trigger Event in effect” branch in the RMBS principal waterfall — an alternate mode gated by a deal-state trigger.

Cross-asset-class observations

What ALL deals have in common

1. **Numbered priority of payments** — the prospectus is always a numbered list of steps.
2. **Each step has multiple targets in a list with explicit order semantics** (sequential, pro-rata, concurrent).
3. **Each step has explicit cap semantics:** “to its planned balance” / “until retired” / “to zero” / “amount equal to...”
4. **Triggers gate behavior**, but the type of trigger varies: stepdown date + cumulative loss + delinquency in RMBS; acceleration / event of default in auto.
5. **A “named amount” abstraction** is always implied — every deal defines named formulas like “Group X Principal Distribution Amount”, “First Priority Principal Payment”, “Net Monthly Excess Cashflow”, and refers to them by name in subsequent steps.
6. **Residual / equity / depositor** is always the final step.

What VARIES across asset classes

Feature	Agency MBS	Non-Agency RMBS	Prime Auto
Distinct collateral groups	YES (1-9 per deal)	YES (typically 1-2)	NO
Separate INT/PRIN sub-streams	YES (INT_CASH/PRIN_CASH)	YES (Interest Remittance Amount + Principal Remittance Amount)	NO (combined Available Funds)
Group-aware allocation	YES	YES (Group 1 Certs, Group 2 Certs)	N/A

Feature	Agency MBS	Non-Agency RMBS	Prime Auto
PAC / TAC / Z behavior	YES	NO	NO
Stepdown date gate	NO	YES	NO
OC / Reserve account in waterfall	RARE (REMIC trust-level fees)	YES (excess interest cascade)	YES (step 8)
Sequential vs pro-rata switch	NO	YES (Sequential Trigger Event)	NO
Interleaved I/P by class	NO	NO (separate I and P waterfalls)	YES (interest 3, prin 4; interest 5, prin 6...)
Loss allocation cascade	rare	YES (reverse seniority)	rare (covered by OC)
Computed (named) distribution amounts	LIGHT (just “Group N Cash Flow Distribution Amount”)	HEAVY (every step references named formulas)	MEDIUM (priority principal payments are formulas)
Recursive splits	YES (FNR 2019-17 Group 7)	RARE	NO
Aggregate Group bond bundles	YES	NO	NO

Cross-cutting needs

- Multi-target rules with explicit payment_style** — already supported in IR. **Use it more** in the FNR fixture and the irGenerator.
- Named distribution amounts** — referenced by name in rule sources. Currently the IR has `INT_CASH` / `PRIN_CASH` as built-in streams and `SPLIT_CASH` to create new ones; this generalizes to “any named formula”. Need a `CalculationNode`-style “ComputedAmount” object that produces a per-period scalar usable as a `from_source` cap.
- if / elif / else over rule blocks** — the RMBS principal waterfall has six mutually exclusive blocks (A through F) keyed on (stepdown date, trigger event). Currently expressed via `condition_trigger` on each rule, which works but is verbose and error-prone (every rule must remember to invert the condition for the “else” branch). A `WaterfallBranch` node with ordered `if / elif / else` cases reads exactly like the prospectus.
- Bond-level loss treatment** — the existing `PAY_WRITEDOWN` rule already does the reverse-seniority cascade. What’s missing is the bond’s *response* to a writedown: does the bond’s balance decrease (`WRITEDOWN`) or stay constant with a deferred carryover (`NOTIONAL_HOLD`)? Critical for CRT (writedown is the primary economic event), important for non-agency RMBS subs.
- Aggregate Group abstraction** — a NAMED collection of bonds treated as a unit (own planned balance, internal allocation rule). Currently expressed by tagging each bond and listing them in a rule’s `to_targets`; works but loses the “this is a group” intent.
- Reserve account integration** — already supported via `PAY_TO_RESERVE` / `PAY_FROM_RESERVE_*`. Used heavily in auto; present in non-agency RMBS via OC mechanics.
- Excess interest sub-cascade** — a separate sub-waterfall driven by what’s left after the “main” interest cascade. Need a clean way to express the carry-forward.

Gaps in the current IR

After this 4-deal sample:

Gap	Severity	Affects
Named distribution amounts (computed scalars) referenced by source name	HIGH	non-agency RMBS, prime auto
Multi-target sequential rules under-used in fixtures and irGenerator	HIGH	every asset class (visual / authoring)
Aggregate Group bond bundles	MEDIUM	agency MBS
Recursive / nested splits	MEDIUM	agency MBS (FNR 2019-17 pattern)
Bond-level loss treatment (writedown vs notional-hold)	HIGH	CRT, non-agency RMBS, future CMBS
Conditional waterfall blocks (<code>if / elif / else</code> over rule groups)	MEDIUM	non-agency RMBS, CRT, auto post-acceleration
Net WAC reserve / sub-cascade plumbing	LOW (covered by <code>SPLIT_CASH</code> + accounts)	non-agency RMBS
Post-acceleration alternate waterfall	LOW (covered by triggers)	auto
Capped fee with overflow to later step	LOW	auto

Proposed IR additions (DRAFT — not approved)

A. Multi-target rule consolidation (no schema change, fixture rewrite + synth/irGen change)

Stop fragmenting by emitting one rule per bond. Both the FNR fixture authors and the `emitPacTacSchedule` block generator should emit multi-target rules (`to_targets: [PA, PB, PC, PD, EO]`, `payment_style: SEQUENTIAL`, `cap_mode: PLANNED`).

Why first: zero schema risk, biggest immediate benefit. The rule count drops from ~24 to ~5 for FNR Group 1.

B. ComputedAmount node (small schema add)

```
class ComputedAmountNode(BaseModel):
    name: str          # "Group 1 Principal Distribution Amount"
    expression: str     # "principal + advance_prin + recovery"
    description: str
```

Rules can reference `ComputedAmountNode` names in `from_sources` the same way they reference `INT_CASH` / `PRIN_CASH`. The runtime resolves the name to the per-period scalar.

C. WaterfallBranch — if / elif / else over rule blocks

The natural way to express “if Trigger Event in effect: pay rules 1; else: pay rules 2” is exactly that — an explicit `if / elif / else` chain at the IR level, not per-rule `condition_trigger` tags. Per-rule tagging works mathematically but is error-prone (you have to remember to invert the condition on every “else” rule and keep the inversion in sync) and unreadable (the prospectus phrasing “If X, do Y, else do Z” is fragmented across many rules).

```
class WaterfallBranch(BaseModel):
    branch_id: str
    description: str
    cases: list[WaterfallCase] # ordered; first matching case fires

class WaterfallCase(BaseModel):
    when: str | None          # trigger name, None for the `else` arm
    invert: bool = False      # treat the trigger as `not when`
    expr: str | None          # OR a free expression evaluated per period
    rules: list[RuleNode]    # rules to execute when this case matches
```

Reads exactly like the prospectus:

```
waterfall_rules:
- branch_id: principal_waterfall
  description: "RMBS principal allocation: stepdown × trigger event"
  cases:
    - when: TriggerEvent
      rules: [ ... rules block A: sequential to senior, no mezz ... ]
    - when: StepdownDate
      invert: true
      rules: [ ... block B: pre-stepdown sequential ... ]
    - expr: "stepdown_date_reached and not trigger_event"
      rules: [ ... block C: post-stepdown pro-rata ... ]
    - when: null # the `else` arm
      rules: [ ... default block ... ]
```

Why prefer this over per-rule `condition_trigger`:

1. **Mutual exclusion is explicit.** `if/elif/else` semantics guarantee exactly one branch fires. Per-rule tags can accidentally fire two branches if conditions aren't perfectly complementary.
2. **DRY.** Don't repeat the same condition on 14 rules.
3. **Reads like the prospectus.** “Block A — Trigger Event in effect” maps to one case; “Block B — pre-stepdown” to another.
4. **Refactor-friendly.** Adding a rule to a conditional block is one append; under the per-rule scheme it's an append plus matching the condition exactly.

When to still use per-rule `condition_trigger`: one-off gates on a *single* rule (e.g., “this single fee only applies if the servicer is in default”). Per-rule remains for one-off cases; `WaterfallBranch` is for multi-rule blocks.

This is the IR equivalent of asking “why don't we just write `if/else`?” — the answer is “we should, and that's what this node is.”

D. AggregateGroup bond bundle

```
class AggregateGroupDef(BaseModel):
    name: str          # "Aggregate Group III"
```

```
members: list[str] # ["MA", "ME", "MC", "MG", "MB"]
schedule_contract: list[...] | None # planned balance (PAC)
internal_allocation: PaymentStyle # SEQUENTIAL / PRO_RATA
```

A virtual bundle that rules can target by name. The runtime caps at the bundle's planned balance and distributes internally per the allocation rule.

E. Bond-level loss treatment (the real loss-allocation question)

Reverse-seniority loss allocation isn't a special new IR construct — it's the standard pattern across RMBS, CRT, and future CMBS, and the existing `PAY_WRITEDOWN` rule type already expresses it:

```
rule_type: PAY_WRITEDOWN
from_sources: [LOSS]
to_targets: [B-2, B-1, M-2, M-1] # reverse seniority
payment_style: SEQUENTIAL
```

The **real question** that the IR currently doesn't answer is: **when a loss hits a bond, what happens to that bond's balance and accrual base going forward?** Two distinct treatments exist in real prospectuses:

1. **WRITEDOWN** — bond balance is reduced by the loss amount. Future coupon accrues on the *reduced* balance. Future principal distributions are based on the reduced balance. This is the CRT default and the non-agency RMBS subordinate default.
2. **NOTIONAL_HOLD** — bond balance stays unchanged. The loss becomes a deferred-amount carryover that reduces cash to that bond until covered by recovery / excess interest, but the bond continues to accrue coupon on its *full* original balance. This is rare but appears in some prime jumbo deals and in particularly investor-friendly subordinate structures.
3. **NONE** — bond doesn't absorb losses at all (typically senior-most class with full guarantee).

Adding this is a small `BondDef` extension:

```
class LossTreatment(StrEnum):
    WRITEDOWN = "WRITEDOWN"
    NOTIONAL_HOLD = "NOTIONAL_HOLD"
    NONE = "NONE"

class BondDef(BaseModel):
    ...
    loss_treatment: LossTreatment = LossTreatment.NONE
    writeup_enabled: bool = False # CRT-style loss reversal
```

The runtime change is small: in `PAY_WRITEDOWN`, look up each target bond's `loss_treatment` and either decrement its balance (`WRITEDOWN`) or accumulate a deferred amount (`NOTIONAL_HOLD`). Recovery / writeup applies inversely.

Why this matters for CRT. CRT bonds are pure notional-bearing instruments — there is no actual cash collateral in the trust. The bonds' economic existence IS their balance, and loss allocation IS the primary cashflow event. Without `loss_treatment` on `BondDef`, the IR can't model CRT at all. The existing `PAY_WRITEDOWN` rule type is the *cascade*; the missing piece is the *bond's* response to that cascade.

F. Recursive SPLIT_CASH (no schema change, runtime change)

Already supported in principle: a `SPLIT_CASH` target stream can be the source of another `SPLIT_CASH`. Verify this works for FNR 2019-17 Group 7's nested 16.67 / 83.33 / first / second pattern.

Round 3 schema review — separating bond identity from rule behavior

This round addresses an architectural concern surfaced during review: the current schema **conflates three different concepts** on `BondDef` and in the rule type enum:

1. **What a bond IS** (intrinsic identity that doesn't change with rules) — IO vs PO vs cash-pay vs Z vs residual vs pseudo
2. **What a bond HAS** (intrinsic schedules / properties) — coupon, notional, schedule contract, maturity, tracking relationships, loss treatment
3. **HOW a bond gets paid** (extrinsic, lives on the rule that pays it) — sequential vs pro-rata, cap mode, conditional gating

The dividing line: if the property is true regardless of which waterfall is paying the bond, it belongs on `BondDef`. If it describes a particular rule's allocation behavior, it belongs on the rule. The current schema crosses this line in several places.

G. Collapse TrancheType + TrancheBehavior into a single TrancheKind

Current state. `TrancheType` has 13 values (`SEQUENTIAL`, `PAC`, `PAC_II`, `TAC`, `SUPPORT`, `Z_BOND`, `ACCRETION_DIRECTED`, `FLOATER`, `INVERSE_FLOATER`, `IO`, `PO`, `PSEUDO`, `RESIDUAL`). `TrancheBehavior` has 4 (`SEQUENTIAL`, `PAC`, `TAC`, `Z`). They overlap, and both encode rule-behavior properties on the bond.

The 13 values are mixing five orthogonal concepts:

Concept	Values currently in <code>TrancheType</code>	Belongs on
Cashflow identity	<code>IO</code> , <code>PO</code> , <code>RESIDUAL</code> , <code>PSEUDO</code> , <code>Z_BOND</code>	<code>BondDef</code> (the bond IS this)
Schedule type	<code>PAC</code> , <code>PAC_II</code> , <code>TAC</code>	<code>BondDef.schedule_contract</code> (already has <code>schedule_type</code>)
Coupon style	<code>FLOATER</code> , <code>INVERSE_FLOATER</code>	<code>BondDef.coupon_type</code> (already exists)
Payment role	<code>SUPPORT</code>	derived from <code>support_tranches</code> relationships
Allocation order	<code>SEQUENTIAL</code> , <code>ACCRETION_DIRECTED</code>	rule's <code>payment_style</code> + bond's accretion targets

Proposed. A single `TrancheKind` that answers only “what *kind* of bond is this”:

```
class TrancheKind(StrEnum):
    CASH_PAY = "CASH_PAY" # ordinary bond – pays cash interest + cash principal
    PAC = "PAC" # planned amortization class – REQUIRES schedule_contract
    TAC = "TAC" # targeted amortization class – REQUIRES schedule_contract
    IO = "IO" # interest-only, notional-bearing
    PO = "PO" # principal-only, zero coupon
    Z = "Z" # accrues during accretion phase, then pays cash
    RESIDUAL = "RESIDUAL" # equity / sweep
    PSEUDO = "PSEUDO" # accounting sink (fees, residuals tracking quantities)
```

8 kinds. **Validation rule:** if `kind ∈ {PAC, TAC}` then `schedule_contract` MUST be non-empty; if `kind ∈ {CASH_PAY, IO, PO, Z, RESIDUAL, PSEUDO}` then `schedule_contract` MUST be empty (or absent). The bond’s identity drives the schedule requirement, not the other way around — Pydantic enforces this with a model validator.

Why keep PAC and TAC as kinds (not just “CASH_PAY with a schedule”):

- **The prospectus uses these names.** “Class PA is a PAC Class. Class TA is a TAC Class.” The IR vocabulary should match.
- **They are distinct economic identities.** A PAC has a *band* of speeds it tolerates; a TAC has a *single* target speed and no upside protection. Both concepts have to be captured somewhere; making them kinds is more legible than inferring from `schedule_speed_low == schedule_speed_high`.
- **Validation enforcement.** A PAC bond without a schedule is always wrong. The kind makes that constraint explicit and catchable at IR validation time, not at runtime.

Everything else moves out:

- **SUPPORT** → derived: a bond is “support” if some other bond’s `relations` list points to it as `SUPPORTED_BY`. The bond itself is just a `CASH_PAY`. Removing the SUPPORT enum value doesn’t lose information.
- **PAC_II** → still a `PAC` kind. The “II” is a structural ordering relative to `PAC_I` (`PAC_I` shortfalls cascade to `PAC_II`’s planned balance), expressible via the rule’s order and a relation pointing `PAC_II` at `PAC_I`. Doesn’t need to be a separate kind.
- **FLOATER / INVERSE_FLOATER** → already lives in `coupon_type` enum (`FIXED | FLOATING | INVERSE_FLOATING | ZERO`). Duplicate.
- **SEQUENTIAL** → not a property of the bond. It’s the rule’s `payment_style`.
- **ACCRETION_DIRECTED** → encoded by Z’s `ACCRETES_TO` relation (see proposal H below).

TrancheBehavior is fully redundant with the simplified schema and should be deleted.

H. Unify tranche relationships into one structured list

Current state. Three different fields express bond-to-bond structural relationships:

- `support_tranches: list[str]` — PAC’s support stack (the bonds that absorb prepay variability for me)
- `supported_by_tranches: list[str]` — Z’s accretion targets (where my PIK accrual is directed while they’re outstanding)
- `parent_tranche: str | None + relation_type: StructureRelation`
 - `notional_ratio: float | None + tracks_bonds: dict[...]` — a tangled mix used for IO/PO notional tracking and inverse-floater parenting

`StructureRelation` has only 3 values (`floaters_inverse`, `io_po`, `z_accrual`), which doesn’t cover the real-world set of relationships you raised: POs, IOs, **inverse IOs**, **inverse floaters**, **super floaters**, MACR exchanges.

Proposed. One `tranche_relations` list with a typed enum that covers every real-world relationship:

```
class TrancheRelationType(StrEnum):
    SUPPORTED_BY = "SUPPORTED_BY" # PAC -> support stack
    ACCRETES_TO = "ACCRETES_TO" # Z -> bonds receiving Z's PIK
    NOTIONAL_TRACKS = "NOTIONAL_TRACKS" # IO / inverse IO (notional follows other bond's balance)
    BALANCE_TRACKS = "BALANCE_TRACKS" # PO mirroring another bond's principal
```

```

COUPON_INVERSE_OF = "COUPON_INVERSE_OF" # inverse floater pegged to a floater
COUPON_LEVERAGE_OF = "COUPON_LEVERAGE_OF" # super floater (multiple of an index)
MACR_EXCHANGE = "MACR_EXCHANGE" # exchangeable / combinable

```

```

class TrancheRelation(BaseModel):
    relation_type: TrancheRelationType
    targets: list[str] # bonds at the other end of the relationship
    weights: list[float] | None # for tracking aggregate baskets
    leverage: float | None # for COUPON_LEVERAGE_OF (e.g., 2.5x)
    cap: float | None # for COUPON_INVERSE_OF (cap - rate * leverage)
    floor: float | None
    description: str = ""

```

```

class BondDef(BaseModel):
    ...
    relations: list[TrancheRelation] = []

```

This single field covers:

Bond type	relations example
Plain IO	<code>[{relation_type: NOTIONAL_TRACKS, targets: [PA, PB, PC, PD], weights: [...]}]</code>
Inverse IO	Same as IO, plus a separate <code>coupon_type=INVERSE_FLOATING</code>
Plain PO	<code>[{relation_type: BALANCE_TRACKS, targets: [...]}]</code> (or just stand-alone with <code>coupon_type=ZERO</code>)
Inverse Floater	<code>[{relation_type: COUPON_INVERSE_OF, targets: [Floater], cap: 0.10}]</code>
Super Floater	<code>[{relation_type: COUPON_LEVERAGE_OF, targets: [Index], leverage: 2.5, cap: 0.12}]</code>
PAC support stack	PAC has <code>[{SUPPORTED_BY, targets: [WA, WB, ...]}]</code>
Z bond	Z has <code>[{ACCRETES_TO, targets: [TA, TB]}]</code>
MACR / RCR	<code>[{MACR_EXCHANGE, targets: [...], weights: [...]}]</code>

The 4 separate fields collapse into 1, gaining expressiveness for inverse IOs / super floaters / MACR while staying compact.

I. Coupon as a value or a schedule (not just a scalar)

Current state. `BondDef.coupon: float | None` plus separate `cap` / `floor` scalars. Cannot express:

- Step-up coupons (Verus 2024-9 — coupon goes from `c` to `c + 1.00%` at month 60)
- Step-down coupons
- Lockout-period coupons (initial fixed rate, then floats)

Proposed. Allow each rate field to be either a scalar or a period-keyed schedule:

```

RateOrSchedule = float | list[RateScheduleEntry]

class RateScheduleEntry(BaseModel):
    from_period: int # inclusive
    rate: float

class BondDef(BaseModel):
    ...
    coupon: RateOrSchedule | None
    margin: RateOrSchedule | None # floater spread
    cap: RateOrSchedule | None
    floor: RateOrSchedule | None

```

Most bonds keep the simple scalar form. Step-up bonds use:

```

coupon:
- { from_period: 1, rate: 5.50 }
- { from_period: 60, rate: 6.50 }

```

J. Notional, not “size”

Current state. `size_dollars: float | None` and `size_pct: float | None`.

Proposed rename. “Size” is ambiguous. Use `notional`:

```


notional: float | None          # dollar notional / par balance
notional_pct_of_collateral: float | None # 0..1 of pool original balance

```

For IOs the notional is the literal IO notional (no principal flow). For cash-pay bonds the notional is the par/face balance. Same field, semantically clearer.

K. Two-phase schedule derivation — speeds are design-time, schedule_contract is runtime

The right framing: **schedule_contract is a derived artifact**. The user owns the schedule's *inputs* (speed band, prepay model, collateral assumptions); the system computes and caches the *output* (the per-period planned balance vector); the runtime reads only the output.



Derivation algorithm (PAC):

For each period t , compute the projected outstanding balance of the bond at speed_low and speed_high using the prepay model and collateral assumptions. The PAC's planned balance at t is the **maximum** of the two (or, equivalently, the balance under the slower speed) — this is what gives PAC stability across the band: fast-prepay scenarios drain support tranches; slow-prepay scenarios let support tranches absorb the shortfall.

For TAC: same calculation, single speed, balance = projected balance at that speed.

Schema after Round 3:

```

class BondDef(BaseModel):
    ...
    kind: TrancheKind          # PAC | TAC | CASH_PAY | ...

    # Design-time inputs (kept so the schedule is re-derivable)
    schedule_speed_low: float | None          # PAC low end; TAC uses this as the target
    schedule_speed_high: float | None         # PAC high end; TAC sets == low
    schedule_model_type: PrepayModelType      # PSA | CPR | ABS | CUSTOM_VECTOR
    schedule_tolerance_bps: float | None

    # Runtime canonical (what the cap_mode=PLANNED rule consults)
    schedule_contract: list[ScheduleEntry]    # [{period, target_balance}]

```

The redundant `schedule_speed_target` field is dropped — `schedule_speed_low == schedule_speed_high` already encodes TAC as the degenerate band.

When does schedule_contract get re-derived?

- User changes the speed band in the structuring UI.
- User changes prepay model or collateral assumptions.
- User changes the bond's notional (changes the absolute balance curve).
- User adjusts which support tranches absorb shortfalls (changes the schedule shape under stress).

The structuring UI runs derivation eagerly on input change; the runtime never derives — it only consumes. This separation gives us:

- **Reproducibility:** a deal definition with a populated `schedule_contract` runs identically every time, regardless of whether speeds / collateral assumptions changed *after* the schedule was built. The IR is self-contained.
- **Auditability:** the `schedule_contract` is what was *actually used* in pricing and execution. Speeds are inputs; the schedule is the audit artifact.
- **Performance:** derivation only happens at design time; the runtime does not pay derivation cost on every period.

- **UI ergonomics:** the user works in speed-band terms (“100%- 250% PSA”), which is how prospectuses and analysts talk; derivation handles the math.

Validation contract:

- If `kind ∈ {PAC, TAC}` and `schedule_contract` is empty → invalid IR (can’t run a PAC bond without a schedule).
- If `kind = PAC`, both `schedule_speed_low` and `schedule_speed_high` should be set.
- If `kind = TAC`, both should be set and equal.
- The validator can additionally check that `schedule_contract` is *consistent* with the speeds and collateral (re-derive and compare with tolerance). This is an optional integrity check, not a correctness requirement at runtime.

L. Drop CONCURRENT from PaymentStyle

`CONCURRENT` is a synonym for `PRO_RATA`. Drop it.

```
class PaymentStyle(StrEnum):
    SEQUENTIAL = "SEQUENTIAL"
    PRO_RATA   = "PRO_RATA"
```

M. Reserve rules are not a separate rule type; they’re rules sourced from / targeted to accounts

Current state. Five reserve-related rule types:

- `PAY_TO_RESERVE` — deposit into reserve
- `PAY_FROM_RESERVE` — generic withdrawal
- `PAY_FROM_RESERVE_INTEREST` — withdrawal labeled “interest”
- `PAY_FROM_RESERVE_PRINCIPAL` — withdrawal labeled “principal”
- (plus the same shape for recourse: `PAY_RECOURSE_INTEREST`, `PAY_RECOURSE_PRINCIPAL`)

This conflates **two orthogonal things**: what cash is moving (interest? principal?) and where it’s coming from / going to (account? bond? stream?). Every “from reserve” rule is really “`PAY_INTEREST` (or `PAY_PRINCIPAL`) with `from_sources` set to a reserve account.”

Proposed. Three rule types covering all account interactions:

```
PAY_INTEREST      # to: bond. from: any stream OR any account.
PAY_PRINCIPAL     # to: bond. from: any stream OR any account.
PAY_TO_ACCOUNT    # to: account. from: any stream OR any account.
                  # (renamed from PAY_TO_RESERVE – accounts are not just reserves)
```

The current `PAY_FROM_RESERVE_INTEREST` becomes `PAY_INTEREST` with `from_sources=[ReserveAcct]`. The current `PAY_TO_RESERVE` becomes `PAY_TO_ACCOUNT`. Recourse is just a named source stream (`from_sources=[RECOURSE_LINE]`).

This change makes accounts truly first-class. Any rule can deposit to one or withdraw from one.

Risk. The runtime *might* attach extra semantics to the reserve-specific rule types (e.g., maintaining a “shortfall carryover” counter that’s affected only by `PAY_FROM_RESERVE_INTEREST`). That semantic is worth preserving but should be moved to the rule’s effect, not its type — e.g., a `tracks_carryover_for: str` field on the rule that names which bond’s interest-shortfall ledger to update. This needs runtime verification before code change.

N. Drop alias tokens; rename INT_CASH / PRIN_CASH

Current built-in tokens (some redundant):

Token	Status
<code>CASH</code>	keep (canonical pool cashflow)
<code>COLLATERAL</code>	drop — alias for <code>CASH</code> , doesn’t add anything
<code>INT_CASH</code>	rename to <code>CASH_INT</code> for prefix consistency
<code>PRIN_CASH</code>	rename to <code>CASH_PRIN</code> for prefix consistency
<code>LOSS</code>	keep
<code>GROUP_<id>_CASH</code>	keep
<code>GROUP_<id>_COLLATERAL</code>	drop — alias
<code>GROUP_<id>_INT_CASH</code>	rename to <code>GROUP_<id>_CASH_INT</code>
<code>GROUP_<id>_PRIN_CASH</code>	rename to <code>GROUP_<id>_CASH_PRIN</code>
<code>GROUP_<id>_LOSS</code>	keep

The `CASH_*` prefix pattern reads better and groups the cashflow streams as siblings under a parent `CASH` concept.

O. AccountType is a label, not runtime semantics

Current `AccountType` enum: `RESERVE | PREFUNDING | REVOLVING | PAYMENT | SPREAD_ACCOUNT`.

(The earlier IR reference in this doc listed the wrong values — `CUSTODIAL | DISTRIBUTION` — and has been corrected.)

Verified against the runtime: `account_type` is purely a display label. The runtime touches it in exactly two places, both of which are passthrough:

1. **Initialization** (`runtime.py:324`) — copy `account_def.account_type.value` (a string) into `AccountWorkspace.account_type`.
2. **Output** (`runtime.py:1611`) — copy that string into the `DealAccountRow` for output / reporting.

There is **no** `if account_type == X` branch anywhere in the runtime. The runtime treats every account uniformly: a named bucket with `balance`, `deposit`, `withdrawal`, and a `required_minimum` array.

MinimumBasis (after the Round 3 Q fix) **IS** the actual behavioral driver for account floors and starting amounts — the runtime now branches on it correctly. `AccountType` remains a passthrough display label.

This means real-world account variety:

- **Reserve account** — credit / liquidity reserve
- **Prefunding account** — holds cash before bonds buy collateral
- **Revolving account** — master-trust style revolving funding
- **Payment / collection account** — temporary holding
- **Spread account** — excess spread tracking
- **Capitalized interest account** — capitalized interest during prefunding
- **Yield supplement account** — auto YSOC
- **Trustee fee reserve** — fee carve-out

...is supported today not by adding more `AccountType` values, but by configuring the right `minimum_basis`, the right `starting_amount`, and the right rules that deposit to / withdraw from the account. For example:

Real-world account	How it's expressed today
Reserve floored at 0.5% of original collateral	<code>minimum_basis: COLLATERAL_BALANCE</code> , <code>minimum_pct: 0.5</code> , plus a <code>PAY_TO_RESERVE</code> rule
Prefunding account drawing down to zero	<code>starting_amount: <prefunded \$></code> , <code>minimum: 0</code> , plus a deposit rule funded by the prefunding stream and a withdrawal rule that pays it out as principal
Capitalized interest account auto-amortizing	<code>starting_amount: <cap-i \$></code> , <code>minimum: 0</code> , plus a <code>PAY_INTEREST</code> rule with this account as <code>from_sources</code> for the bond's coupon during the cap-i period
YSOC account	A <code>SPREAD_ACCOUNT</code> (label) feeding a <code>SPLIT_CASH</code> that boosts under-WAC bond cash

Proposed change: rename `AccountType` → `AccountCategory` to reflect that the field is a UI / reporting label and stop signaling “this might gate behavior” via the name. No runtime change needed.

If a future deal needs runtime behavior keyed on account type (e.g., “PREFUNDING accounts auto-amortize without explicit rules”), that should be added as a *separate* explicit field — e.g., `auto_amortizes: bool`, `amortization_schedule: list[...]` — not by overloading `account_type` with hidden semantics.

Q. Implement `minimum_basis` and `starting_basis` per-period semantics — FIXED

Status: Implemented. `runtime.py` now honors all 4 basis values for both `minimum_basis` and `starting_basis`. New regression tests in `tests/test_account_minimum_basis.py` lock in the per-period behavior. The remainder of this section is preserved for context.

Original state (silently buggy). The IR had `MinimumBasis` and `starting_basis` fields with 4 enum values:

```
class MinimumBasis(StrEnum):
    FIXED_DOLLAR = "FIXED_DOLLAR"
    COLLATERAL_BALANCE = "COLLATERAL_BALANCE"
    NOTE_BALANCE = "NOTE_BALANCE"
    ORIGINAL_COLLATERAL = "ORIGINAL_COLLATERAL"
```

Intended semantics (per the field name and the per-period `required_minimum` array shape):

Value	Intended behavior
<code>FIXED_DOLLAR</code>	<code>floor = minimum_amount</code> (constant \$)

Value	Intended behavior
COLLATERAL_BALANCE	$\text{floor}[t] = \text{minimum_pct} \times \text{pool_balance}[t]$ (steps down as pool amortizes)
NOTE_BALANCE	$\text{floor}[t] = \text{minimum_pct} \times \text{outstanding_note_balance}[t]$ (steps down as bonds amortize)
ORIGINAL_COLLATERAL	$\text{floor} = \text{minimum_pct} \times \text{original_pool_balance}$ (constant)

Actual runtime behavior (`runtime.py:310-321`, `_allocate_account_workspace`):

```
minimum = float(account_def.minimum_amount or 0.0)
if account_def.minimum_pct is not None:
    minimum = max(minimum, collateral_balance_0 * float(account_def.minimum_pct) / 100.0)
required_minimum = np.zeros(cf_len)
required_minimum[:] = minimum
```

The runtime:

1. Computes the minimum **once** at initialization
2. Always uses `collateral_balance_0` (the *original* balance) as the percentage basis
3. Broadcasts the single value across all periods
4. **Never reads `minimum_basis` to decide how to compute**

So `minimum_basis = COLLATERAL_BALANCE` and `minimum_basis = ORIGINAL_COLLATERAL` produce identical runtime behavior. Same for `NOTE_BALANCE` (silently treated as `ORIGINAL_COLLATERAL`).

The `starting_basis` field has the same problem — `starting_pct` always uses `collateral_balance_0`, ignoring the basis enum.

Practical impact. A reserve account authored with `minimum_basis=COLLATERAL_BALANCE` and `minimum_pct=0.5%` expecting the floor to amortize down with the pool will instead have the floor fixed at $0.5\% \times \text{original_balance}$ for the deal's entire life. For long-amortizing pools this is a meaningful overstatement of the reserve floor late in the deal's life and can mask reserve breaches.

Proposed fix. Implement per-period recomputation by honoring the enum value. The fix lives at the start of each period's account-evaluation pass (or hoisted into a derived array if the basis value depends only on bond/pool state):

```
def _period_account_minimum(
    account_def: AccountDef,
    period: int,
    pool_balance_t: float,
    note_balance_t: float,
    collateral_balance_0: float,
) -> float:
    pct = account_def.minimum_pct or 0.0
    floor_pct = {
        MinimumBasis.FIXED_DOLLAR: 0.0,
        MinimumBasis.COLLATERAL_BALANCE: pool_balance_t * pct / 100.0,
        MinimumBasis.NOTE_BALANCE: note_balance_t * pct / 100.0,
        MinimumBasis.ORIGINAL_COLLATERAL: collateral_balance_0 * pct / 100.0,
    }[account_def.minimum_basis]
    return max(account_def.minimum_amount or 0.0, floor_pct)
```

Same shape for `starting_basis` at period 0 only.

Why this matters. Real-world reserve mechanics step down with amortization in many deals — auto ABS reserves often have a “step-down floor” that's “0.50% of current pool balance with a \$X minimum”; non-agency RMBS reserves track note balance through the OC waterfall. Currently neither is correctly modeled.

Severity. This was a HIGH-severity silent correctness bug because the IR accepted the value, the validator passed, the runtime silently ignored the user's choice, and the output looked plausible but was wrong for any account with a non-`ORIGINAL_COLLATERAL` basis.

Resolution (May 2026). Fixed in `runtime.py`. The new `_allocate_account_workspace` builds `required_minimum` per period based on `minimum_basis`:

- `FIXED_DOLLAR` — broadcast `minimum_amount` to every period.
- `ORIGINAL_COLLATERAL` — broadcast `max(minimum_amount, minimum_pct × collateral_balance_0)`.
- `COLLATERAL_BALANCE` — vectorized `max(minimum_amount, minimum_pct × pool_balance[t])` per period using the projected pool balance series.

- `NOTE_BALANCE` — period 0 from the initial note stack at allocation time; periods 1..T are filled by `_refresh_note_balance_minimums(deal, accounts, bonds, t)` inside the main run loop, which sums `bonds[*].balance[t]` for `is_bond=True, is_pseudo=False` bonds and applies `minimum_pct`.

`starting_basis` is honored at period 0 in the same way.

Test coverage in `tests/test_account_minimum_basis.py`:

- `FIXED_DOLLAR` floor stays constant; `pct` is ignored.
- `ORIGINAL_COLLATERAL` floor stays constant at `pct × initial_pool`.
- `COLLATERAL_BALANCE` floor monotonically decreases with pool amortization.
- `NOTE_BALANCE` floor monotonically decreases with bond amortization and tracks the live note stack.
- `minimum_amount` (dollar floor) takes precedence when it exceeds the percentage-derived floor.
- `starting_basis` produces the right initial balance for every enum value.

P. Schema cleanup migration table

Current	Proposed	Rationale
<code>TrancheType</code> (13 values)	<code>TrancheKind</code> (8 values: CASH_PAY, PAC, TAC, IO, PO, Z, RESIDUAL, PSEUDO)	Drop SEQUENTIAL, SUPPORT, ACCRETION_DIRECTED, PAC_II, FLOATER, INVERSE_FLOATER (rule behavior, derived role, or coupon-type properties); KEEP PAC and TAC because they're real economic identities and validation can enforce that PAC/TAC require <code>schedule_contract</code>
<code>schedule_speed_target</code> field	(deleted)	Redundant — TAC uses <code>schedule_speed_low == schedule_speed_high</code> as the degenerate band
<code>TrancheBehavior</code> (4 values)	(deleted)	Fully redundant with <code>TrancheKind</code> + <code>schedule_contract</code> + <code>pay_mode</code>
<code>support_tranches</code> , <code>supported_by_tranches</code> , <code>parent_tranche</code> , <code>relation_type</code> , <code>notional_ratio</code> , <code>tracks_bonds</code>	<code>relations: list[TrancheRelation]</code>	One typed list covers PAC support, Z accretion, IO/PO tracking, inverse floater, super floater, MACR
<code>coupon: float</code> , <code>cap: float</code> , <code>floor: float</code> , <code>margin: float</code>	<code>RateOrSchedule</code> (scalar OR period-keyed schedule)	Step-up / lockout / time-conditional coupons
<code>size_dollars</code> , <code>size_pct</code>	<code>notional</code> , <code>notional_pct_of_collateral</code>	Naming clarity (covers IO notional and cash-pay par)
<code>schedule_speed_low</code> , <code>schedule_speed_high</code> , <code>schedule_speed_target</code>	<code>schedule_speed_band: tuple[float, float]</code> (or move to <code>schedule_metadata</code>)	TAC = degenerate PAC band
<code>PaymentStyle.CONCURRENT</code>	(deleted)	Synonym of <code>PRO_RATA</code>
<code>RuleType.PAY_FROM_RESERVE</code> , <code>PAY_FROM_RESERVE_INTEREST</code> , <code>PAY_FROM_RESERVE_PRINCIPAL</code> , <code>PAY_RECOURSE_*</code>	(deleted)	Use <code>PAY_INTEREST</code> / <code>PAY_PRINCIPAL</code> with account or recourse stream as <code>from_sources</code>
<code>RuleType.PAY_TO_RESERVE</code>	<code>PAY_TO_ACCOUNT</code>	Accounts are not just reserves
Built-in token <code>COLLATERAL</code>	(deleted)	Alias for <code>CASH</code>
Built-in token <code>GROUP_<id>_COLLATERAL</code>	(deleted)	Alias for <code>GROUP_<id>_CASH</code>
<code>INT_CASH</code> / <code>PRIN_CASH</code> (and group variants)	<code>CASH_INT</code> / <code>CASH_PRIN</code> (and group variants)	Prefix consistency
<code>AccountType</code> enum (treated as runtime-significant)	<code>AccountCategory</code> (UI label only)	Verified: runtime never branches on <code>account_type</code>
<code>minimum_basis</code> and <code>starting_basis</code> ignored by runtime	Per-period recomputation honoring the enum value	FIXED (May 2026) — proposal Q. <code>_allocate_account_workspace</code> + <code>_refresh_note_balance_minimums</code> in <code>runtime.py</code> ; covered by <code>tests/test_account_minimum_basis.py</code> (9 tests, all passing).

Net schema impact:

- 2 enums collapsed (`TrancheType` 13→8 as `TrancheKind`, `TrancheBehavior` deleted)
- 6 fields on `BondDef` collapsed into 1 list (`relations`)
- 1 `BondDef` field deleted (`schedule_speed_target`)
- 5 rule types deleted (4 reserve, 1 recourse — folded into the generic interest/principal rules)

- 1 rule type renamed (`PAY_TO_RESERVE` → `PAY_TO_ACCOUNT`)
- 1 enum value deleted (`PaymentStyle.CONCURRENT`)
- 4 built-in tokens cleaned up (drop 2 aliases; rename 2 + group variants)
- 1 enum repurposed (`AccountType` → `AccountCategory` UI label)

The result is a schema where:

- A bond is one of 8 things (kinds), with PAC/TAC carrying a validator that requires a `schedule_contract` .
- Schedules / coupons / notionals / relationships live on the bond as concrete data, not behavior tags.
- Schedule derivation is design-time; the runtime consumes the derived `schedule_contract` directly without re-deriving.
- Behavior — sequential vs pro-rata, schedule-cap vs cleanup, account-sourced vs cash-sourced — lives entirely on the rule.

Open questions for user before any code change

1. **Priority.** Which of (A)–(F) above matter for your near-term deal universe? My read: (A) is must-have *now* (it fixes the visual problem you flagged); (B) and (C) are needed once we start modeling private-label RMBS; (D) is agency-MBS-specific;
 1. is **must-have for CRT** and important for non-agency RMBS subordinates; (F) is already kind of working.
2. **More research breadth.** Do you want me to extend this study to the other asset classes (subprime auto, credit card, CLO, marketplace, equipment, solar) before proposing changes, or are the structural patterns above enough to converge on the abstractions for now and revisit the IR when those asset classes are needed? **My recommendation: do (A) now (no IR risk), park (B)–(F) until we have ≥5 deals per relevant asset class.**
3. **The PAC/TAC/Z question** you raised originally. The IR *intentionally* does not have `PAY_PRINCIPAL_PAC_SCHEDULE` as a distinct rule type — the PAC behavior is encoded as `BondDef.tranche_behavior=PAC` + `BondDef.schedule_contract` + `RuleNode.cap_mode=PLANNED` . This is the right call for the *runtime* (PAC and SEQ have the same allocation logic; only the bond’s per-period cap differs). But for the *UI* we should recognize “this rule pays a PAC bond with `cap_mode=PLANNED`” and render it via the existing `pay_pac_schedule` Blockly block. The synthesizer can detect this from the IR (any rule whose targets include a bond with `tranche_behavior=PAC` and whose `cap_mode=PLANNED` is a PAC schedule rule).
The IR doesn’t need a new rule type for PAC/TAC; the synthesizer + irGenerator just need to recognize the pattern.

Additional deals (round 2 research)

Ginnie Mae REMIC 2025-203 (single-family, multi-group)

Three Security Groups; Group 3 has the canonical PAC + Z + Support structure verbatim:

SECURITY GROUP 3 — *The Group 3 Principal Distribution Amount and the Accrual Amount will be allocated in the following order of priority: 1. Sequentially, to BP and PB, in that order, until reduced to their Aggregate Scheduled Principal Balance for that Distribution Date 2. To Z, until retired 3. Sequentially, to BP and PB, in that order, without regard to their Aggregate Scheduled Principal Balance, until retired*

Identical structure to FNR. “Aggregate Scheduled Principal Balance” is the same abstraction as Fannie’s “Aggregate Group Planned Balance”. Confirms the FNR pattern is industry-standard for agency REMICs.

Ginnie Mae REMIC 2025-009 (HECM-backed reverse mortgage)

Reverse-mortgage REMIC has a distinct structural feature — the **Deferred Interest Amount**. The 3-step waterfall is:

1. Concurrently, to AI, FA and FB, **pro rata based on their respective Interest Accrual Amounts**, up to the Class Interest Accrual Amount
2. Concurrently, to FA and FB, **pro rata based on their Class Principal Balances**, in reduction of their Class Principal Balances, up to the Principal Distribution Amount
3. To AI, until the **Class AI Deferred Interest Amount** is reduced to zero

Step 3 is unique — it pays accrued-but-unpaid IO interest only *after* the principal cascade. Standard MBS pays IO at the same time as bond interest. **Reverse-mortgage REMICs need a “deferred interest catch-up” rule type or a `cap_mode=DEFERRED` flag.** Not in current IR.

Ginnie Mae REMIC 2024-115 (Multifamily)

*Allocation of Principal: On each Distribution Date, a percentage of the Principal Distribution Amount will be applied to the **Trustee Fee**, and the remainder of the Principal Distribution Amount (the “**Adjusted Principal Distribution Amount**”) will be allocated, sequentially, to A and B, in that order, until retired.*

Multifamily REMICs **deduct trustee fees from principal** as a percentage *before* the cascade. Single-family agency REMICs don't do this (trustee/servicer fees are netted at the underlying MBS layer). The IR's `PAY_FEE` rule with `basis_type=COLLATERAL_BALANCE` covers this; the multifamily case puts the fee earlier in the cascade.

Freddie Mac REMIC general structure (offering circular)

Freddie Mac REMICs introduce a structural concept that doesn't exist in our current IR: **Single-Tier vs Double-Tier Series**:

- **Single-Tier**: REMIC Certificates represent direct beneficial ownership in *one* REMIC Pool. Cashflows flow pool → classes.
- **Double-Tier**: An **Upper-Tier REMIC Pool** + one or more **Lower-Tier REMIC Pools** stacked. Lower-Tier Classes are themselves **Mortgage Securities** of the Upper-Tier Pool. Cashflows flow pool → Lower-Tier classes → Upper-Tier classes.

This is **REMIC-inside-a-REMIC** structure. Most agency deals abstract this away (the cashflow analyst just looks at the ultimate classes). For our IR, we already model this implicitly — the `Loan` → `BMA cashflow engine` → `DealRunInput` chain treats the underlying agency MBS as a black box delivering a single pool's cashflows to the deal. The lower-tier internal mechanics are out of scope unless we want to model RCR / MACR exchanges.

Also notable: **MACR (Modifiable and Combinable REMIC) Certificates** allow exchange of one class for proportionate interests in another. This is purely an issuance/secondary mechanic (no waterfall implication) and doesn't touch the IR.

Verus Securitization Trust 2024-9 (Modern Non-QM RMBS)

Adds three patterns not in JPMMT 2006:

1. **Step-up coupon at year 5**. "On each payment date beginning in January 2029, the A1, A2 and A3 [classes] will receive the sum of [the deal's] fixed coupon, plus 1.00%."
 - Time-conditional coupon change. Requires the IR to support coupon as a function of period, not a static field.
 - The current `BondDef.coupon: float` does not capture this. `BondDef.coupon_schedule: list[{period, rate}]` would.
2. **Last Cashflow (LCF) class**. "A1A, A1B, A1LCF" — A1LCF gets paid LAST among the seniors. New seniority sub-pattern.
 - LCF works as a normal bond with a junior position within the senior tier; expressible via the existing rule ordering.
3. **Class XS** — explicit excess-spread tracking class. Class XS noteholders also control optional-redemption rights.
 - Excess-spread tracking is just a residual class with extra economic rights; expressible via `tranche_type=RESIDUAL`.
 - Optional-redemption is governance, not waterfall, and is out of scope for the IR.

Pay structure: "senior notes pro rata, mezz + sub sequentially." Same hybrid pattern as JPMMT 2006 post-stepdown, but encoded with no stepdown date — Non-QM deals typically don't have one because the loans amortize quickly.

Toyota Auto Receivables 2024-A (Prime Auto, Toyota)

Confirms the Ford Credit prime auto shape. Differences:

- Has **First Priority Principal Distribution Amount** (Class A catch-up) AND **Second Priority Principal Distribution Amount** (Class A+B catch-up) AND **Regular Principal Distribution Amount** (target OC build) — same three named amounts.
- **"Yield Supplement Overcollateralization Amount"** (YSOC) — a parallel concept used in adjusting the "adjusted pool balance" for sub-WAC loans. Specific to auto deals where some loans have rates below the WAC needed to support the bonds. Pure computation tweak, doesn't change waterfall structure.
- Reserve account replenishment, capped trustee fees with overflow — same as Ford / SDART.
- Confirms prime auto = same grammar across sponsors.

Toyota Lexus Owner Trust 2024-A (Auto LEASE — different shape)

Lease-backed ABS has a SIMPLER waterfall (8 steps, not 11-14):

1. Servicing Fee
2. Transaction Fees and Expenses (capped \$300K/yr)
3. Note Interest (pro rata across all classes — NOT class-by-class)
4. Note Principal — priority principal distribution amount (single combined amount, not per-class)
5. Reserve Account Deposit
6. Note Principal — regular principal distribution amount
7. Additional Transaction Fees and Expenses (overflow)
8. Excess Amounts to certificateholder

Differences vs auto loan ABS:

- **Pro-rata interest across ALL classes** (no class-by-class step). Less common in loan ABS.
- **One combined "priority principal distribution amount"** instead of per-class first/second priority. The lease structure pays all notes pro rata for principal, with no class-priority cascade.

- “**Securitization Value**” instead of “Pool Balance” — the lease collateral is valued differently because of residual risk on the vehicles.
- Otherwise structurally the same as auto loan ABS.

The IR already supports this (multi-target PAY_INTEREST with PRO_RATA payment_style; multi-target PAY_PRINCIPAL with PRO_RATA). Just simpler.

Westlake Automobile Receivables Trust 2024-1 (Subprime Auto)

8 classes of notes (A-1, A-2, A-3 + B, C, D, E). Same waterfall shape as SDART — interleaved I/P, named priority allocations, reserve replenishment as a step. Westlake confirms the subprime auto pattern is consistent across sponsors.

Updated cross-asset matrix

After 13 deals across 5 asset classes:

Feature	Agency MBS	Agency CRT	Non-Agency RMBS	Prime Auto	Subprime Auto	Auto Lease
Distinct collateral groups	YES (1-9)	NO (single ref pool)	YES (1-2)	NO	NO	NO
Separate INT/PRIN sub-streams	YES	YES (synthetic)	YES	NO	NO	NO
PAC / TAC / Z behavior	YES	NO	NO	NO	NO	NO
Stepdown date gate	NO	YES	YES	NO	NO	NO
Reserve account in waterfall	RARE	NO	YES	YES	YES	YES
Sequential vs pro-rata switch	NO	YES (gate on perf trigger)	YES	NO	NO	NO
Interleaved I/P by class	NO	NO	NO	YES	YES	NO (pro-rata)
Bond writedown on losses	NO	YES (core mechanic)	YES	NO	NO	NO
Bond writeup on loss reversal	NO	YES (rare)	RARE	NO	NO	NO
Named computed distribution amounts	LIGHT	MEDIUM	HEAVY	MEDIUM	MEDIUM	LIGHT
Recursive splits	YES	NO	RARE	NO	NO	NO
Aggregate Group bond bundles	YES	NO	NO	NO	NO	NO
Step-up coupon	RARE	NO	YES (Non-QM)	NO	NO	NO
Deferred interest catch-up	RARE (HECM)	NO	NO	NO	NO	NO
Acceleration alternate waterfall	NO	NO	NO	YES	YES	YES

Updated IR gap list

Adding the new patterns from round 2:

Gap	Severity	Affects	Already in IR?
Multi-target rule consolidation (authoring)	HIGH	every class	NO (fixture/irGen issue, not IR)
Named computed distribution amounts	HIGH	non-agency RMBS, auto	NO
minimum_basis and starting_basis is honored at runtime	HIGH → FIXED (May 2026)	every account-using deal	YES — implemented in <code>runtime.py</code> ; regression tests in <code>tests/test_account_minimum_basis.py</code> (proposal Q)
Aggregate Group bond bundles	MEDIUM	agency MBS	NO
Recursive SPLIT_CASH	MEDIUM	agency MBS	PARTIAL (need runtime verify)

Gap	Severity	Affects	Already in IR?
Bond-level loss treatment (writedown vs notional_hold)	HIGH	CRT (core), non-agency RMBS	NO
Conditional waterfall blocks (if / elif / else)	MEDIUM	non-agency RMBS, CRT, auto post-accel	PARTIAL (per-rule trigger only)
Bond writeup on loss reversal	LOW	CRT	NO
Time-conditional bond coupons (step-ups)	MEDIUM	Non-QM, callable seniors	NO
Acceleration alternate waterfall	MEDIUM	every auto deal	PARTIAL (trigger-based)
Deferred interest catch-up (HECM IO)	LOW	reverse mortgage REMIC	NO
Net WAC reserve sub-cascade	LOW	non-agency RMBS	PARTIAL (SPLIT_CASH + accounts)
Capped fee with overflow to later step	LOW	auto	PARTIAL (max_amount_fixed)

Part 2 — IR reference: the schema and how to write a deal

This second half of the document covers the IR itself. Goal: make the IR both **highly reusable** (one schema for many asset classes) and **human-readable** (an analyst should be able to read the IR top-to-bottom and match it to the prospectus paragraph by paragraph).

IR overview

The IR is a JSON-serializable Pydantic schema that captures one deal's entire static structure in a single document. It is:

- **Source of truth.** The runtime executes the IR; the UI (Blockly canvas + Properties panel) edits the IR; saved deals persist as IR JSON; tests compare against the IR.
- **Versioned.** `schema_version: "1.0.0"` allows forward-compatible migrations.
- **Self-validating.** Pydantic enforces field types; the top-level `_validate_references` validator enforces cross-field consistency (every rule's `from_sources` and `to_targets` must reference declared bonds, accounts, fees, or built-in streams; every rule's `condition_trigger` must reference a declared trigger; every bond's `support_tranches` must reference real bonds; etc.).
- **Round-tripping.** IR → runtime → cashflow output is deterministic; IR → Blockly synthesizer → workspace → `irGenerator` → IR is meant to be lossless.

Top-level schema: `DealDefinition`

```
class DealDefinition(BaseModel):
    schema_version: str = "1.0.0"
    deal_name: str                                     # "FNR 2006-018 (Group 1 + Group 2)"
    description: str = ""
    origination_date: date | None
    settlement_date: date | None

    # Structure
    bonds: list[BondDef]                               # All tranches, including pseudo bonds
    accounts: list[AccountDef]                         # Reserves, prefunding, custodial
    fees: list[FeeDef]                                 # Trustee, servicer, etc.
    triggers: list[TriggerNode]                       # Conditional gates
    calculations: list[CalculationNode]               # Named expressions for triggers
    waterfall_rules: list[RuleNode]                   # The actual priority of payments
    collateral_groups: list[CollateralGroupDef]        # Multi-pool deals

    # Solver / overrides / runtime extensions (see below)
    deal_knobs: dict[str, Any] = {}
```

`deal_knobs` — what's actually in there

`deal_knobs` is intentionally a free-form `dict[str, Any]` because it serves four distinct purposes. **Five reserved keys** have specific runtime meaning; everything else is a free scalar that gets injected into expression-evaluation context.

Key	Type	Purpose
<code>source_formulas</code>	<code>dict[str, str]</code>	Named per-period expressions that become first-class <code>from_source</code> tokens. Example: <code>{"NET_EXCESS": "INT_CASH - bond_coupon_total"}</code> makes <code>NET_EXCESS</code> referenceable in any rule's <code>from_sources</code> .
<code>balance_trackers</code>	<code>dict[str, str]</code>	Maps a residual / pseudo bond's name to a runtime quantity it should mirror. Example: <code>{"R": "collateral_balance"}</code> makes residual <code>R</code> carry a balance equal to outstanding pool balance each period. Used when a residual interest's economics track a non-bond quantity.
<code>orig_collat_bal_override</code>	<code>float</code>	Override for the original collateral balance used in trigger denominator calculations (when the natural pool balance differs from the trigger reference balance, e.g., tape excludes some loans).
<code>allow_negative_cashflow_math</code>	<code>bool</code>	Permit transient negative intermediate cash states during rule execution. Used for deals with negative-amortization streams (HECM, certain ARM cases) where cash conservation invariants need temporary relaxation.

Key	Type	Purpose
<any_identifier>	int \ float	Free expression-context globals. Every numeric value with an identifier-safe key gets injected into the <code>expr</code> evaluation context for fee <code>amount_expr</code> , rule <code>max_amount_expr</code> , calculation <code>expression</code> , and trigger thresholds. Example: <code>{"servicing_fee_bps": 25.0}</code> lets a fee expression reference <code>servicing_fee_bps / 10000</code> directly.

The four functional roles, in plain language:

1. **Solver scratch-pad.** The solver writes proposed values into `deal_knobs.<name>` and re-runs the deal. `KnobSpec.knob_path` uses dot-paths like `deal_knobs.class_a_pctbal` to reach them.
2. **Expression-context globals.** Numeric values are auto-injected into the runtime expression evaluator so any expression-bearing field (`amount_expr`, `max_amount_expr`, `condition_expr`, trigger thresholds, calculation expressions) can reference them by bare name without declaring them as `CalculationNode`s.
3. **Runtime feature flags.** A few well-known boolean keys toggle runtime behavior (currently just `allow_negative_cashflow_math`).
4. **Runtime extension hooks.** `source_formulas` and `balance_trackers` are reserved-key escape hatches for capabilities that haven't been promoted to first-class IR fields yet. Both are candidates for elevation: `source_formulas` should become the proposed `ComputedAmountNode`; `balance_trackers` should become a `BondDef` field (`tracks_quantity: str | None`).

Migration intent for `deal_knobs`

Per the human-readability principle “no hidden runtime knobs,” `deal_knobs` is a tension we accept temporarily. The direction of travel is to **shrink** the schemaless surface:

- `source_formulas` → graduate to `calculations: list[ComputedAmountNode]` (proposed addition B).
- `balance_trackers` → graduate to `BondDef.tracks_quantity` field.
- `orig_collat_bal_override` → graduate to a top-level `DealDefinition.original_collateral_balance: float | None`.
- `allow_negative_cashflow_math` → graduate to a `DealDefinition.runtime_options: RuntimeOptions` typed nested model (one field per flag).

After graduation, `deal_knobs` becomes purely #1 + #2: a typed scratch-pad of named scalar overrides for the solver and for expression evaluation, with everything behavioral promoted to real schema fields.

BondDef — every tranche the deal pays

A bond is anything that receives cash from the waterfall. This includes residual classes and “pseudo bonds” used as accounting sinks for fees.

```
class BondDef(BaseModel):
    name: str # "PA", "Class A-1", "M-1"
    # NOTE: tranche_type currently has 13 values mixing identity, schedule, coupon, and role.
    # Round 3 proposal G collapses this to TrancheKind (8 values:
    # CASH_PAY | PAC | TAC | IO | PO | Z | RESIDUAL | PSEUDO). PAC and TAC remain
    # because they're real economic identities; the validator enforces that
    # PAC/TAC kinds require a non-empty schedule_contract. SUPPORT becomes a
    # derived role; FLOATER/INVERSE_FLOATER live in coupon_type; SEQUENTIAL is
    # rule behavior, not bond identity; PAC_II is just a PAC with structural ordering.
    tranche_type: TrancheType # SEQUENTIAL | PAC | PAC_II | TAC | SUPPORT | Z_BOND |
    ACCRETION_DIRECTED | FLOATER | INVERSE_FLOATER | IO | PO | PSEUDO | RESIDUAL
    # NOTE: tranche_behavior is fully redundant with the simplified TrancheKind + schedule + pay_mode.
    # Round 3 proposal G deletes this field entirely.
    tranche_behavior: TrancheBehavior # SEQUENTIAL | PAC | TAC | Z
    is_bond: bool # False for residual/pseudo
    is_pseudo: bool # True for fee sinks

    # Coupon — Round 3 I: each rate field should accept a scalar OR a period-keyed
    # schedule (RateScheduleEntry list) so step-up / lockout coupons can be expressed.
    coupon_type: CouponType # FIXED | FLOATING | INVERSE_FLOATING | ZERO
    coupon: float | None # Annual percent (5.5 = 5.5%)
    margin: float | None # Floater spread over index
    index_name: str | None # SOFR | TERM_SOFR_1M | etc.
    cap: float | None # Floater rate cap
    floor: float | None # Floater rate floor
```

```

# Sizing – Round 3 J: rename to `notional` and `notional_pct_of_collateral`.
# "Notional" is the universal term (covers IOs which have notional but no principal flow)
# and "size" is ambiguous.
size_dollars: float | None          # → notional
size_pct: float | None              # → notional_pct_of_collateral (0..1, not 0..100)

# Maturity / accrual
maturity_date: date | None
day_count: DayCount                 # 30/360 | ACT/360 | ACT/365 | ACT/ACT
accrual_period: AccrualPeriod       # MONTHLY | QUARTERLY | SEMI_ANNUAL | ANNUAL

# PAC / TAC parameters
# Round 3 K: two-phase derivation – the speed fields are DESIGN-TIME inputs
# (kept so the schedule can be re-derived if collateral / band assumptions
# change). The runtime ONLY consumes `schedule_contract`. Derivation runs
# at structuring-time on input change, caches the result here, and the
# runtime reads the cache directly. TAC is the degenerate band where
# low == high; `schedule_speed_target` is therefore redundant and is dropped.
schedule_type: ScheduleType | None # PAC | TAC | SUPPORT (Round 3 G: drops SUPPORT – the kind tells
you)
schedule_model_type: PrepayModelType # PSA | CPR | ABS | CUSTOM_VECTOR (design-time input)
schedule_speed_low: float | None      # PAC lower PSA / TAC target (design-time input)
schedule_speed_high: float | None     # PAC upper PSA (TAC: == low) (design-time input)
schedule_speed_target: float | None   # Round 3 K: dropped – redundant with low when TAC
schedule_contract: list[dict]         # [{period, target_balance}] – runtime canonical, derived
schedule_tolerance_bps: float | None

# Tranche relationships – Round 3 H: collapse the next 6 fields into ONE
# typed list `relations: list[TrancheRelation]` covering SUPPORTED_BY,
# ACCRETES_TO, NOTIONAL_TRACKS, BALANCE_TRACKS, COUPON_INVERSE_OF,
# COUPON_LEVERAGE_OF, MACR_EXCHANGE – full coverage of POs, IOs, inverse IOs,
# inverse floaters, super floaters, MACR exchange classes.
support_tranches: list[str]           # PAC support stack → relations[?].SUPPORTED_BY
supported_by_tranches: list[str]      # Z accretion targets → relations[?].ACCRETES_TO
parent_tranche: str | None             # IO/PO parent for tracking → relations[?].NOTIONAL_TRACKS /
BALANCE_TRACKS
relation_type: StructureRelation | None # FLOATER_INVERSE | IO_PO | Z_ACCRUAL – Round 3 H expands to 7
values
notional_ratio: float | None          # IO notional ratio → relations[?].weights / leverage
tracks_bonds: dict[str, list[str]] | None # legacy IO/PO tracking → relations[?].targets

# Z-bond / accrual
z_accrual_enabled: bool
z_release_trigger: str | None
accrual_start_period: int | None
accrual_end_period: int | None

# Multi-group
group_id: str | None                  # "GROUP_1", "GROUP_2"

# Solver knob flags
solver_knob_coupon: bool
solver_knob_size: bool

pay_mode: PayMode                     # CASH_PAY | PIK

```

Reusability principle: one `BondDef` covers all bond types. Today the differentiation is split across `tranche_type` + `tranche_behavior` + schedule + accrual + tracking fields — that works but conflates intrinsic identity with rule-driven behavior. The Round 3 cleanup (proposals G + H) reduces this to one `TrancheKind` enum (the bond's identity), one `relations` list (its structural ties to other bonds), and the actual schedule / coupon / notional / accrual data. PAC-ness, TAC-ness, support status, and IO/PO tracking become consequences of those concrete fields — no separate behavior tag needed.

RuleNode — one waterfall step

```

class RuleNode(BaseModel):
    rule_id: str                # "r_int_PA_pacI"
    rule_type: RuleType         # PAY_INTEREST | PAY_PRINCIPAL | ...
    order: int                  # 0, 1, 2 ... priority

```

```

from_sources: list[str]           # ["CASH_INT"] or ["GROUP_1_CASH_PRIN"] or ["ReserveAcct"]
to_targets: list[str]            # ["PA", "PB", "PC", "PD", "EO"] (bonds, accounts, named streams)
payment_style: PaymentStyle     # SEQUENTIAL | PRO_RATA (Round 3: CONCURRENT alias dropped)
cap_mode: CapMode | None        # PLANNED | SCHEDULED | TARGETED | NONE

max_amount_fixed: float | None  # Hard $ cap on this rule
max_amount_expr: str | None     # Computed cap expression
target_weights: list[float] | None # SPLIT_CASH per-target weights (rule-level, not bond-level)

condition_trigger: str | None   # Trigger name
condition_invert: bool         # Run when trigger is FALSE
condition_expr: str | None     # Custom condition expression

reserve_account: str | None     # PAY_TO_RESERVE / PAY_FROM_RESERVE_* (Round 3 M: redundant once
rules can source/target accounts directly)
allow_negative_source: bool

group_id: str | None            # Multi-group routing
description: str = ""

```

Reusability principle: every prospectus step maps to ONE `RuleNode`. The shape of the step (PAC schedule, sequential pay, pro-rata, fee, residual sweep, conditional pay) is conveyed by the combination of `rule_type` + `payment_style` + `cap_mode` + `condition_trigger`. **No new rule type for “PAY_PRINCIPAL_PAC_SCHEDULE”** or similar — those are just `PAY_PRINCIPAL` with `cap_mode=PLANNED` on a bond that has a `schedule_contract` (Round 3 G: the bond’s PAC-ness is the schedule itself, not a separate `tranche_behavior` flag).

RuleType enum — what cash this rule moves

RuleType	Cash moved	Typical sources	Typical targets
<code>PAY_INTEREST</code>	Bond cash interest	typically <code>CASH_INT</code> or <code>CASH</code> ; can be any account or stream	One or more bonds
<code>PAY_INTEREST_SHORTFALL</code>	Catch-up of unpaid interest	typically <code>CASH_INT</code> ; can be any account	Bonds with unpaid coupon
<code>PAY_PRINCIPAL</code>	Principal	typically <code>CASH_PRIN</code> or <code>CASH</code> ; can be any account or stream	One or more bonds
<code>PAY_WRITEDOWN</code>	Loss allocation	<code>LOSS</code>	Bonds (reverse seniority)
<code>PAY_FEE</code>	Fee	any account or stream	One fee payee
<code>PAY_TO_RESERVE</code>	Deposit to an account	any source	Any account
<code>PAY_FROM_RESERVE_INTEREST</code>	Withdraw from account, used for interest	An account	Bonds (interest shortfall)
<code>PAY_FROM_RESERVE_PRINCIPAL</code>	Withdraw from account, used for principal	An account	Bonds (principal acceleration)
<code>PAY_FROM_RESERVE</code>	Generic withdrawal from account	An account	Bonds or other targets
<code>PAY_RECOURSE_INTEREST</code>	Sponsor recourse covering interest	Recourse stream	Bonds (interest shortfall)
<code>PAY_RECOURSE_PRINCIPAL</code>	Sponsor recourse covering principal	Recourse stream	Bonds (principal acceleration)
<code>PAY_RESIDUAL</code>	Residual sweep	Any leftover stream	Residual classes
<code>SPLIT_CASH</code>	Stream plumbing	One source stream	N named virtual streams

(Round 3 cleanup, proposal M: the five reserve-specific rule types (`PAY_TO_RESERVE`, `PAY_FROM_RESERVE`, `PAY_FROM_RESERVE_INTEREST`, `PAY_FROM_RESERVE_PRINCIPAL`) and the recourse rule types (`PAY_RECOURSE_INTEREST`, `PAY_RECOURSE_PRINCIPAL`) conflate “what cash is moving” with “where it’s coming from / going to.” The proposed cleaner model: - `PAY_INTEREST` / `PAY_PRINCIPAL` accept any account or stream as their `from_sources`. So “pay interest from reserve” becomes `PAY_INTEREST from=[ReserveAcct]`. “Pay interest from sponsor recourse” becomes `PAY_INTEREST from=[RECOURSE_LINE]`. - `PAY_TO_RESERVE` is renamed to `PAY_TO_ACCOUNT` (accounts are not just reserves) and accepts any account as `to_targets`. - The “this rule covers an interest shortfall” semantic is moved from the rule type into a small `tracks_carryover_for: str` field on the rule, naming the bond whose shortfall ledger gets decremented.)

PaymentStyle — order semantics within a multi-target rule

Style	Behavior
<code>SEQUENTIAL</code>	Pay first target until its cap, then next, etc. (“In that order”)

Style	Behavior
PRO_RATA	Pay all targets simultaneously by their balance / face / coupon weight

(Round 3 cleanup: `CONCURRENT` was a synonym of `PRO_RATA` and has been removed; older code referencing it should migrate to `PRO_RATA`.)

CapMode — schedule cap interpretation for PAY_PRINCIPAL

Mode	Stops paying when	Prospectus phrasing
PLANNED	Bond reaches its <code>schedule_contract</code> planned balance	“to its Planned Balance”
SCHEDULED	Bond reaches its scheduled (next-period) balance	“to its Scheduled Balance”
TARGETED	Bond reaches a target	“to its Targeted Balance”
NONE	Never (cleanup rule)	“without regard to its Planned Balance” / “until retired”

Built-in source/target tokens (reusable across asset classes)

Tokens that any rule’s `from_sources` / `to_targets` can reference without explicit declaration:

Token	Meaning
CASH	Combined pool cashflow (interest + principal + recovery)
CASH_INT	Pool interest stream only (separated from principal)
CASH_PRIN	Pool principal stream only
LOSS	Pool loss stream (for writedown rules)
GROUP_<id>_CASH	Combined cashflow for collateral group <id>
GROUP_<id>_CASH_INT	Interest stream for collateral group <id>
GROUP_<id>_CASH_PRIN	Principal stream for collateral group <id>
GROUP_<id>_LOSS	Loss stream for collateral group <id>

When a rule declares `group_id`, the bare tokens (`CASH`, `CASH_INT`, `CASH_PRIN`, `LOSS`) are auto-prefixed with `GROUP_<id>_` at compile time. So a multi-group rule can write `from_sources: ["CASH_INT"]` and have it resolve to the right group automatically.

(Round 3 cleanup notes: - The `COLLATERAL` and `GROUP_<id>_COLLATERAL` aliases for `CASH` have been dropped — they didn’t add expressiveness and duplicated the canonical name. - The previous tokens `INT_CASH` / `PRIN_CASH` (and group variants) have been renamed to `CASH_INT` / `CASH_PRIN` for prefix consistency. Existing IR documents using the old names should be migrated; the validator can accept both names during a transition period.)

Other elements

AccountDef — named cash buckets (reserves, prefunding, etc.)

```
class AccountDef(BaseModel):
    name: str # "Reserve_Account"
    # Current code: RESERVE | PREFUNDING | REVOLVING | PAYMENT | SPREAD_ACCOUNT
    # VERIFIED against runtime.py: account_type is a passthrough display label only.
    # The runtime stores it at init (line 324) and copies it to output (line 1611);
    # there is NO `if account_type == X` branch anywhere. Round 3 0 renames to
    # AccountCategory. Behavior comes from minimum_basis + the rules that touch
    # the account; minimum_basis is now (post-Round 3 Q fix) honored at runtime
    # for both starting and required-minimum calculations.
    account_type: AccountType # RESERVE | PREFUNDING | REVOLVING | PAYMENT | SPREAD_ACCOUNT
    starting_amount: float # $ amount at closing
    starting_pct: float | None # OR % of <basis quantity> at period 0
    starting_basis: MinimumBasis # FIXED_DOLLAR | COLLATERAL_BALANCE | NOTE_BALANCE |
    ORIGINAL_COLLATERAL
    # — honored at runtime as of May 2026 (Round 3 Q fix)
    minimum_amount: float # Dollar floor; combined with the pct-derived floor via max()
    minimum_pct: float | None # OR % of <basis quantity> per minimum_basis
    minimum_basis: MinimumBasis # FIXED_DOLLAR (constant $) | COLLATERAL_BALANCE (steps down with
    pool) |
    # NOTE_BALANCE (steps down with notes) | ORIGINAL_COLLATERAL
```

(constant pct of orig)

– honored at runtime as of May 2026 (Round 3 Q fix)

FeeDef — periodic fee paid to a payee

```
class FeeDef(BaseModel):
    name: str # "TRUSTEE", "SERVICER"
    basis_type: FeeBasisType # FIXED_DOLLAR | COLLATERAL_BALANCE | PER_LOAN
    amount: float # $ amount per period (FIXED_DOLLAR)
    amount_expr: str | None # OR computed expression
    rate: float | None # Annual rate (COLLATERAL_BALANCE)
    rate_expr: str | None
    minimum: float # Floor on the fee
    frequency: FeeFrequency # MONTHLY | QUARTERLY | ANNUAL
    cumulative: bool # Track unpaid fees as carryover
```

TriggerNode — conditional gate

```
class TriggerNode(BaseModel):
    name: str # "CumLossTrigger"
    metric_type: TriggerMetricType # CUMULATIVE_LOSS | CUMULATIVE_DEFAULT | DELINQUENCY_RATE | OC_TEST
    | IC_TEST | CUSTOM
    threshold_value: float | None
    threshold_schedule: list[float] | None # Per-period threshold (e.g., RMBS time-stepped triggers)
    calculation_ref: str | None # Pointer to a CalculationNode for dynamic thresholds
    comparison_ref: str | None # ">" / ">=" / etc.
    cure_periods: int | None # How many clean periods to clear the trigger
```

CalculationNode — named expressions

```
class CalculationNode(BaseModel):
    name: str # "cum_loss_pct"
    expression: str # "sum(loss[0:i+1]) / orig_collat_bal"
    description: str
```

Used by triggers (and in proposed [ComputedAmountNode](#)) to compute per-period scalars. Supports a safe subset of Python: arithmetic, min/max/abs, references to bond/account/pool state.

CollateralGroupDef — multi-pool deals

```
class CollateralGroupDef(BaseModel):
    group_id: str # "GROUP_1"
    label: str # Human-readable
    description: str
```

Activates per-group cash routing in the runtime. When non-empty, every non-pseudo bond and every cashflow-touching rule must declare `group_id`.

Worked example: FNR 2006-018 Group 1, written in IR

The prospectus says (verbatim, paraphrased for brevity):

Group 1 cash flow priority of payments: 1. Pay interest on all cash-paying bonds. 2. Pay principal of Aggregate Group I to its Planned Balance (PA → PB → PC → PD → EO sequential). 3. Pay principal of Aggregate Group II to its Planned Balance (TA → TB sequential). 4. Pay principal to Z to zero. 5. Distribute 95.65% of remaining principal to support sequential (WA → WB → WC → WD → WE → WG); 4.35% to PO. 6. Pay principal of Aggregate Group II to zero (without regard to planned balance). 7. Pay principal of Aggregate Group I to zero (without regard to planned balance). 8. Sweep remaining cash to residual.

Below is the **same waterfall expressed compactly in IR** (one rule per prospectus step). Every rule has a stable `rule_id` that matches the prospectus phrasing so an analyst can search the IR and find the corresponding step:

```
deal_name: "FNR 2006-018 Group 1"
collateral_groups:
  - group_id: GROUP_1
    label: "Group 1 (PAC + Z + Support)"

bonds:
  # Aggregate Group I (PAC I + IO/PO)
```



```

- { name: PA,  tranche_type: PAC,    tranche_behavior: PAC, group_id: GROUP_1,
    coupon: 5.5, size_dollars: 33710000, schedule_contract: [...],
    support_tranches: [WA, WB, WC, WD, WE, WG, PO] }
- { name: PB,  tranche_type: PAC,    tranche_behavior: PAC, group_id: GROUP_1, ... }
- { name: PC,  tranche_type: PAC,    tranche_behavior: PAC, group_id: GROUP_1, ... }
- { name: PD,  tranche_type: PAC,    tranche_behavior: PAC, group_id: GROUP_1, ... }
- { name: EO,  tranche_type: PO,     coupon_type: ZERO,   group_id: GROUP_1, ... }
- { name: EI,  tranche_type: IO,     tracks_bonds: { balance: [PA, PB, PC, PD] }, ... }

# Aggregate Group II (PAC II / accretion-directed)
- { name: TA,  tranche_type: PAC,    tranche_behavior: PAC, group_id: GROUP_1, ... }
- { name: TB,  tranche_type: PAC,    tranche_behavior: PAC, group_id: GROUP_1, ... }

# Z-bond (Aggregate Group II support)
- { name: Z,   tranche_type: Z_BOND, tranche_behavior: Z, pay_mode: PIK,
    group_id: GROUP_1, z_accrual_enabled: true, supported_by_tranches: [TA, TB] }

# Support tranches
- { name: WA,  tranche_type: SUPPORT, group_id: GROUP_1, ... }
- { name: WB,  tranche_type: SUPPORT, group_id: GROUP_1, ... }
- { name: WC,  tranche_type: SUPPORT, group_id: GROUP_1, ... }
- { name: WD,  tranche_type: SUPPORT, group_id: GROUP_1, ... }
- { name: WE,  tranche_type: SUPPORT, group_id: GROUP_1, ... }
- { name: WG,  tranche_type: SUPPORT, group_id: GROUP_1, ... }

# Support PO (4.35% of support cash)
- { name: PO,  tranche_type: PO,     coupon_type: ZERO, group_id: GROUP_1, ... }

# Residual
- { name: R,   tranche_type: RESIDUAL, is_pseudo: true }

waterfall_rules:
# Step 1 – Interest cascade (one rule, all cash-paying bonds in priority order)
- rule_id: r_int_cascade
  rule_type: PAY_INTEREST
  order: 0
  group_id: GROUP_1
  from_sources: [INT_CASH]
  to_targets: [PA, PB, PC, PD, TA, TB, EI, WA, WB, WC, WD, WE, WG]
  payment_style: SEQUENTIAL
  description: "Pay accrued bond coupon. Z is PIK; its coupon is capitalized."

# Step 2 – PAC I to its Planned Balance (one rule, all PAC I bonds in priority)
- rule_id: r_prin_pac_i_planned
  rule_type: PAY_PRINCIPAL
  order: 1
  group_id: GROUP_1
  from_sources: [PRIN_CASH]
  to_targets: [PA, PB, PC, PD, EO]
  payment_style: SEQUENTIAL
  cap_mode: PLANNED
  description: "Aggregate Group I to its Planned Balance"

# Step 3 – PAC II to its Planned Balance
- rule_id: r_prin_pac_ii_planned
  rule_type: PAY_PRINCIPAL
  order: 2
  group_id: GROUP_1
  from_sources: [PRIN_CASH]
  to_targets: [TA, TB]
  payment_style: SEQUENTIAL
  cap_mode: PLANNED
  description: "Aggregate Group II to its Planned Balance"

# Step 4 – Z-bond
- rule_id: r_prin_Z
  rule_type: PAY_PRINCIPAL
  order: 3

```



```

group_id: GROUP_1
from_sources: [PRIN_CASH]
to_targets: [Z]
payment_style: SEQUENTIAL
description: "Z to zero"

# Step 5a – 95.65 / 4.35 face-weighted split
- rule_id: r_supp_split
  rule_type: SPLIT_CASH
  order: 4
  group_id: GROUP_1
  from_sources: [PRIN_CASH]
  to_targets: [WAWG_BUCKET, PO_BUCKET]
  target_weights: [0.956521694276, 0.043478305724]
  description: "Face-weighted split of remaining principal: 95.65% to WA-WG, 4.35% to PO"

# Step 5b – WA-WG cascade (multi-target sequential)
- rule_id: r_pay_wawg
  rule_type: PAY_PRINCIPAL
  order: 5
  group_id: GROUP_1
  from_sources: [WAWG_BUCKET]
  to_targets: [WA, WB, WC, WD, WE, WG]
  payment_style: SEQUENTIAL

# Step 5c – Support PO
- rule_id: r_prin_PO
  rule_type: PAY_PRINCIPAL
  order: 6
  group_id: GROUP_1
  from_sources: [PO_BUCKET]
  to_targets: [PO]

# Step 5d – Sweep leftover bucket cash back into PRIN_CASH
- rule_id: r_supp_sweep_back
  rule_type: SPLIT_CASH
  order: 7
  group_id: GROUP_1
  from_sources: [WAWG_BUCKET, PO_BUCKET]
  to_targets: [PRIN_CASH]
  target_weights: [1.0]

# Step 6 – PAC II cleanup
- rule_id: r_prin_pac_ii_uncapped
  rule_type: PAY_PRINCIPAL
  order: 8
  group_id: GROUP_1
  from_sources: [PRIN_CASH]
  to_targets: [TA, TB]
  payment_style: SEQUENTIAL
  cap_mode: NONE
  description: "Aggregate Group II to zero, without regard to Planned Balance"

# Step 7 – PAC I cleanup
- rule_id: r_prin_pac_i_uncapped
  rule_type: PAY_PRINCIPAL
  order: 9
  group_id: GROUP_1
  from_sources: [PRIN_CASH]
  to_targets: [PA, PB, PC, PD, E0]
  payment_style: SEQUENTIAL
  cap_mode: NONE

# Step 7b – Support cleanup (in case face-weighted split left tail residue)
- rule_id: r_prin_supports_uncapped
  rule_type: PAY_PRINCIPAL
  order: 10
  group_id: GROUP_1

```

```

from_sources: [PRIN_CASH]
to_targets: [WA, WB, WC, WD, WE, WG, PO]
payment_style: SEQUENTIAL
cap_mode: NONE

# Step 8 – Residual
- rule_id: r_resid_int
  rule_type: PAY_RESIDUAL
  order: 11
  group_id: GROUP_1
  from_sources: [INT_CASH]
  to_targets: [R]
  payment_style: SEQUENTIAL

- rule_id: r_resid_prin
  rule_type: PAY_RESIDUAL
  order: 12
  group_id: GROUP_1
  from_sources: [PRIN_CASH]
  to_targets: [R]
  payment_style: SEQUENTIAL

```

Rule count: 13 (one per logical waterfall step), versus the current FNR fixture’s **35 rules** (one per bond per step). The behavior is identical — the runtime walks `to_targets` sequentially with the right `cap_mode` — but the IR is now readable as prospectus paragraphs.

Reusability principles

- One bond schema covers all bond types.** The `BondDef` shape accommodates PAC, TAC, Z, IO, PO, sequential, support, residual, pseudo via the `tranche_type` + `tranche_behavior` + schedule / accrual / tracking fields. **Don’t add a `PACBondDef` or `ZBondDef` subclass** — encode the variation in the existing fields.
- One rule schema covers all rule types.** Most prospectus structural variation is captured by the combination of `rule_type` + `payment_style` + `cap_mode` + `condition_trigger`. PAC / TAC / cleanup / mezz cascades / interest waterfalls / principal waterfalls / fee priorities all fit. **Don’t add `PAY_PRINCIPAL_PAC_SCHEDULE` or similar specialized rule types** — encode the variation in `cap_mode` and the bond’s `tranche_behavior`.
- One rule per prospectus step, not one per bond.** The prospectus says “Pay sequentially to PA, PB, PC, PD, EO until their planned balances.” That is **ONE rule** with `to_targets: [PA, PB, PC, PD, EO]` and `cap_mode: PLANNED`. Do NOT fragment it into 5 rules.
- Built-in tokens are reused everywhere.** `CASH` / `INT_CASH` / `PRIN_CASH` / `LOSS` are universal. Multi-group deals add `GROUP_<id>_*` prefixes which the runtime auto-resolves when `group_id` is set on the rule.
- Asset-class differences are parametric, not structural.** Prime auto vs subprime auto: same grammar, more classes / more named amounts / lower fee cap. Agency MBS vs non-agency RMBS: different feature mix (PAC vs OC/stepdown), but the same primitives.

Human-readability principles

- Stable, descriptive `rule_id`s that match the prospectus.** `r_prin_pac_i_planned` reads as “rule, principal, PAC I, to planned balance” — a person can scan the IR and match it to the prospectus paragraph. Avoid synthetic ids like `r_principal_0001` that lose intent.
- `description` field on every rule.** One sentence, ideally verbatim from the prospectus. The IR should be readable on its own without referring back to the source document.
- Group IR top-level lists by purpose, not alphabetically.** `bonds` ordered by seniority; `waterfall_rules` ordered by priority (matching `order`); `triggers` ordered by deal-state relevance. The IR document reads top-to-bottom as the prospectus reads.
- Comments allowed via `description` even where Pydantic doesn’t normally permit.** Every node has a `description: str` field. Use it. Especially for:
 - Rules: paste the prospectus phrasing verbatim.
 - Bonds: note the rationale for any non-obvious field (e.g., why `support_tranches` is set as it is).
 - Triggers: explain the threshold table and cure logic.
- Compact over verbose where math is identical.** If a rule could be written as one multi-target rule or as N single-target rules, prefer the compact form because it matches the prospectus. The runtime produces identical output either way; the compact form is far easier to read and edit.

6. **Avoid hidden runtime knobs.** If a deal needs an exotic behavior (Z accrual, schedule cap mode, face-weighted split), make it visible in the IR via existing fields (`cap_mode` , `target_weights` , `pay_mode=PIK`). Do not bury it in `deal_knobs` or in code branches in the runtime.
7. **One DealDefinition file per real-world deal.** A single, self-contained, version-controlled JSON / YAML / Python factory. The fixture for FNR 2006-018 is one file: `tests/fixtures/fnr_2006_018/deal_definition.py` . The full combined Group 1 + Group 2 deal is one `DealDefinition` with multiple `collateral_groups` , not two separate files.

Anti-patterns to avoid

Anti-pattern	Better approach
One rule per bond (“ <code>r_int_PA</code> , <code>r_int_PB</code> , ...”)	One rule per waterfall step (<code>r_int_cascade</code> with all bonds)
New <code>RuleType</code> for each prospectus phrase	Existing rule types + <code>cap_mode</code> + <code>payment_style</code>
Subclassing <code>BondDef</code> for PAC / Z / IO	Set <code>tranche_type</code> + <code>tranche_behavior</code> + <code>schedule_contract</code> / <code>tracks_bonds</code>
Ad-hoc behavioral key in <code>deal_knobs</code> (not one of the 5 reserved keys)	Visible IR fields (<code>target_weights</code> , <code>cap_mode</code> , <code>pay_mode</code>) — promote new behavior to a real schema field
Synthetic <code>rule_id</code> like <code>r_001</code>	Descriptive <code>r_prin_pac_i_planned</code>
Empty <code>description</code> field	Verbatim prospectus phrasing
Two <code>DealDefinition</code> files for one prospectus deal with multiple groups	One <code>DealDefinition</code> with multiple <code>collateral_groups</code>
Trigger logic inline in expressions	Named <code>TriggerNode</code> referenced by <code>condition_trigger</code>
Computed amount inline in <code>max_amount_expr</code>	Named <code>CalculationNode</code> (or proposed <code>ComputedAmountNode</code>)